

Towards cache-optimal address allocation:
How slow is your code?

Mario Preishuber

Master's Thesis
Department of Computer Sciences
University of Salzburg

Problem Statement

Given a trace T of load and store instructions

find metrics that characterize the trace performance
for a given cache C and

implement an execution engine for computing their
quantities.

Problem Statement

Given a **trace T** of load and store instructions

find **metrics** that characterize the trace **performance**
for a given cache C and

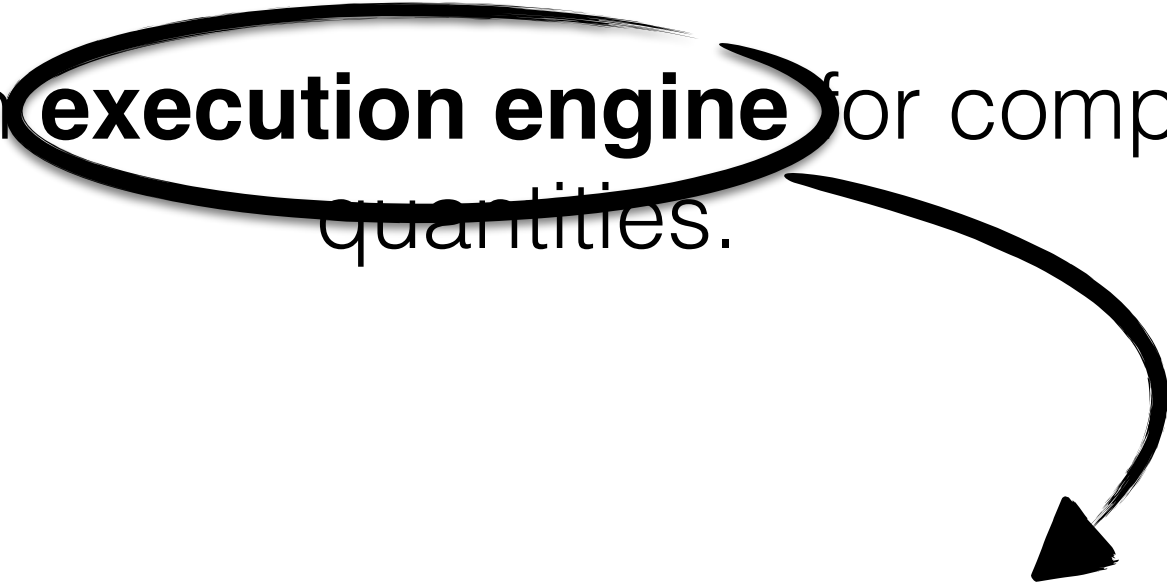
implement an **execution engine** for computing their
quantities.

Problem Statement

Given a **trace T** of load and store instructions

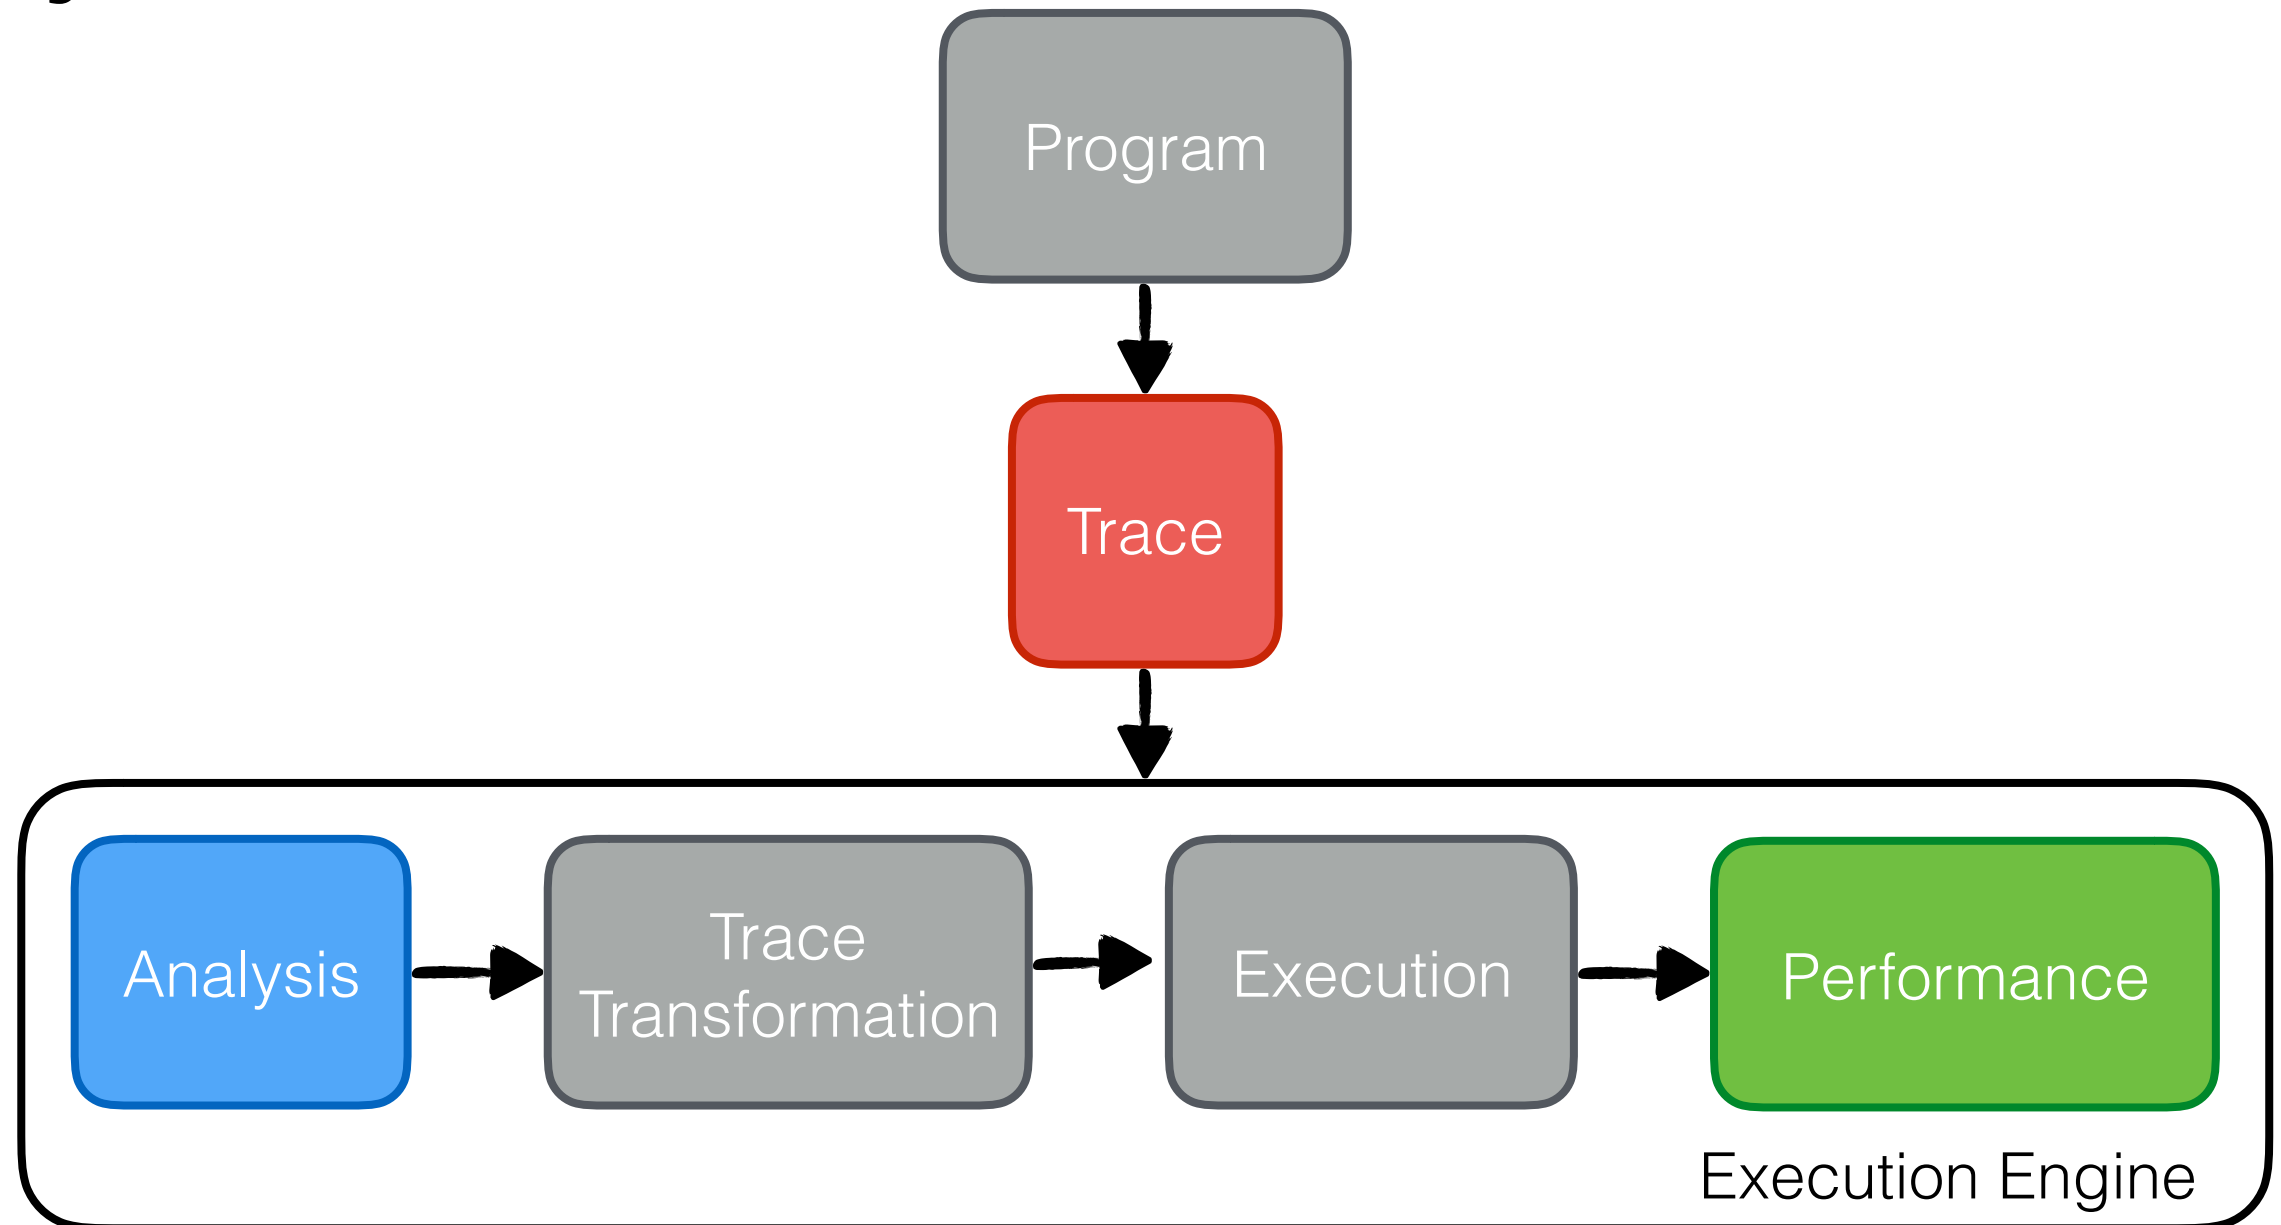
find **metrics** that characterize the trace **performance**
for a given cache C and

implement an **execution engine** for computing their
quantities.



Collaboration with Alexander Miller

Overview System



Preparation

Memory Access Trace

```
void main ()
{
    int sum, x, y, z;
    sum = 0;
    x = 1;
    y = 2;
    sum = sum + x;
    z = 3;
    sum = sum + y;
    sum = sum + z;
}
```

Preparation

Memory Access Trace

```
void main ()
{
    int sum, x, y, z;
    sum = 0;
    x = 1;
    y = 2;
    sum = sum + x;
    z = 3;
    sum = sum + y;
    sum = sum + z;
}
```

```
store &sum
store &x
store &y
load &sum
load &x
store &sum
store &z
load &sum
load &y
store &sum
load &sum
load &z
store &sum
```

Preparation

Memory Access Trace

```
void main ()  
{  
    int sum, x, y, z;  
    sum = 0;  
    x = 1;  
    y = 2;  
    sum = sum + x;  
    z = 3;  
    sum = sum + y;  
    sum = sum + z;  
}
```



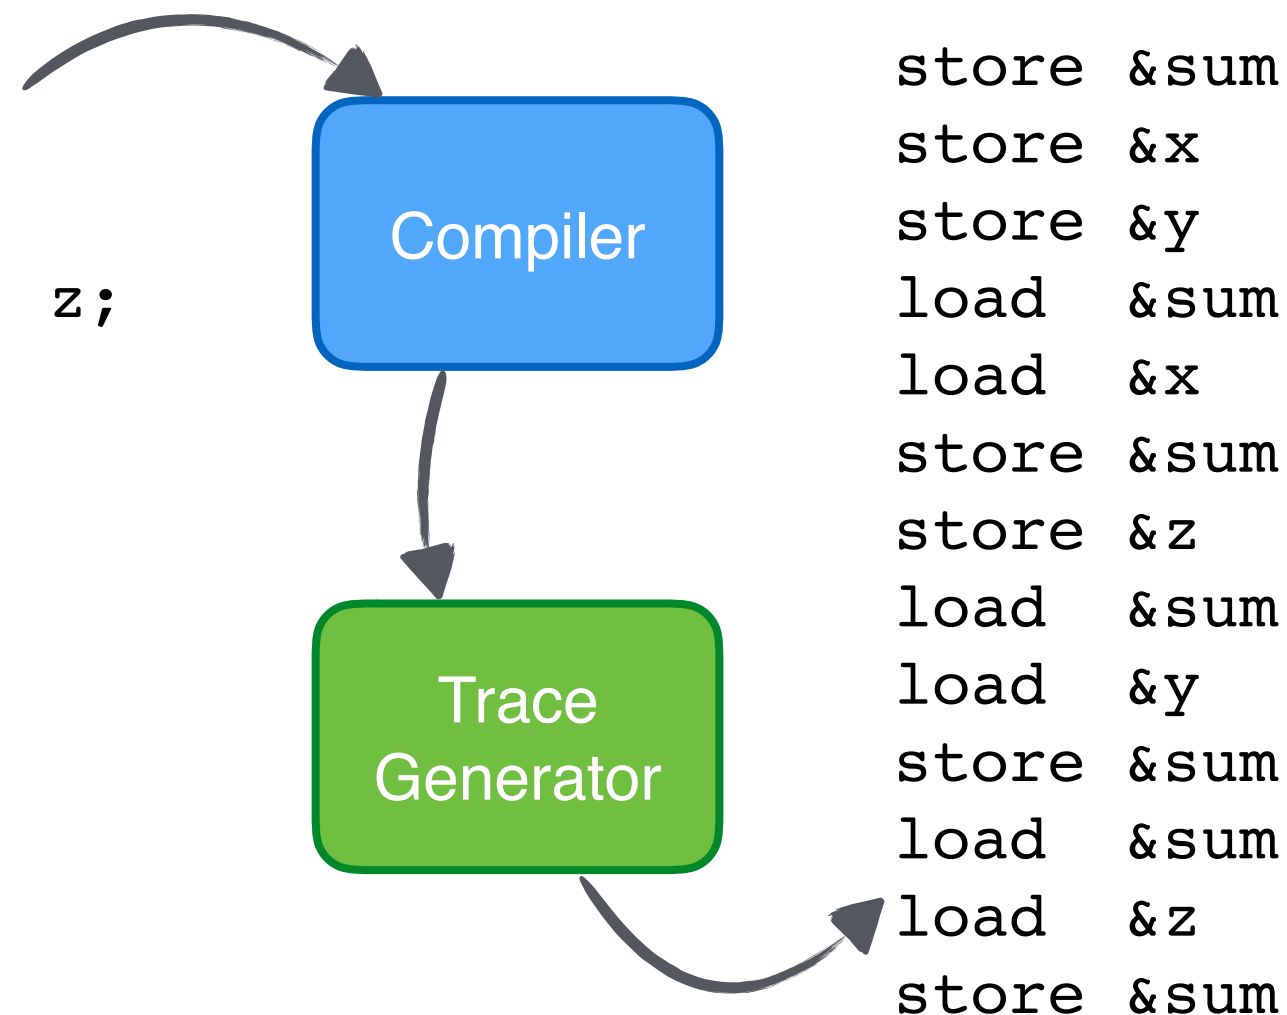
Compiler

```
store &sum  
store &x  
store &y  
load &sum  
load &x  
store &sum  
store &z  
load &sum  
load &y  
store &sum  
load &sum  
load &z  
store &sum
```


Preparation

Memory Access Trace

```
void main ()  
{  
    int sum, x, y, z;  
    sum = 0;  
    x = 1;  
    y = 2;  
    sum = sum + x;  
    z = 3;  
    sum = sum + y;  
    sum = sum + z;  
}
```



Preparation

Memory Access Trace

```
void main ()
{
    int sum, x, y, z;
    sum = 0;
    x = 1;
    y = 2;
    sum = sum + x;
    z = 3;
    sum = sum + y;
    sum = sum + z;
}
```

GCC
v4.8.5

Valgrind
Lackey

store &sum
store &x
store &y
load &sum
load &x
store &sum
store &z
load &sum
load &y
store &sum
load &sum
load &z
store &sum

Preparation

Memory Access Trace

```
void main ()  
{  
    int sum, x, y, z;  
    sum = 0;  
    x = 1;  
    y = 2;  
    sum = sum + x;  
    z = 3;  
    sum = sum + y;  
    sum = sum + z;  
}
```

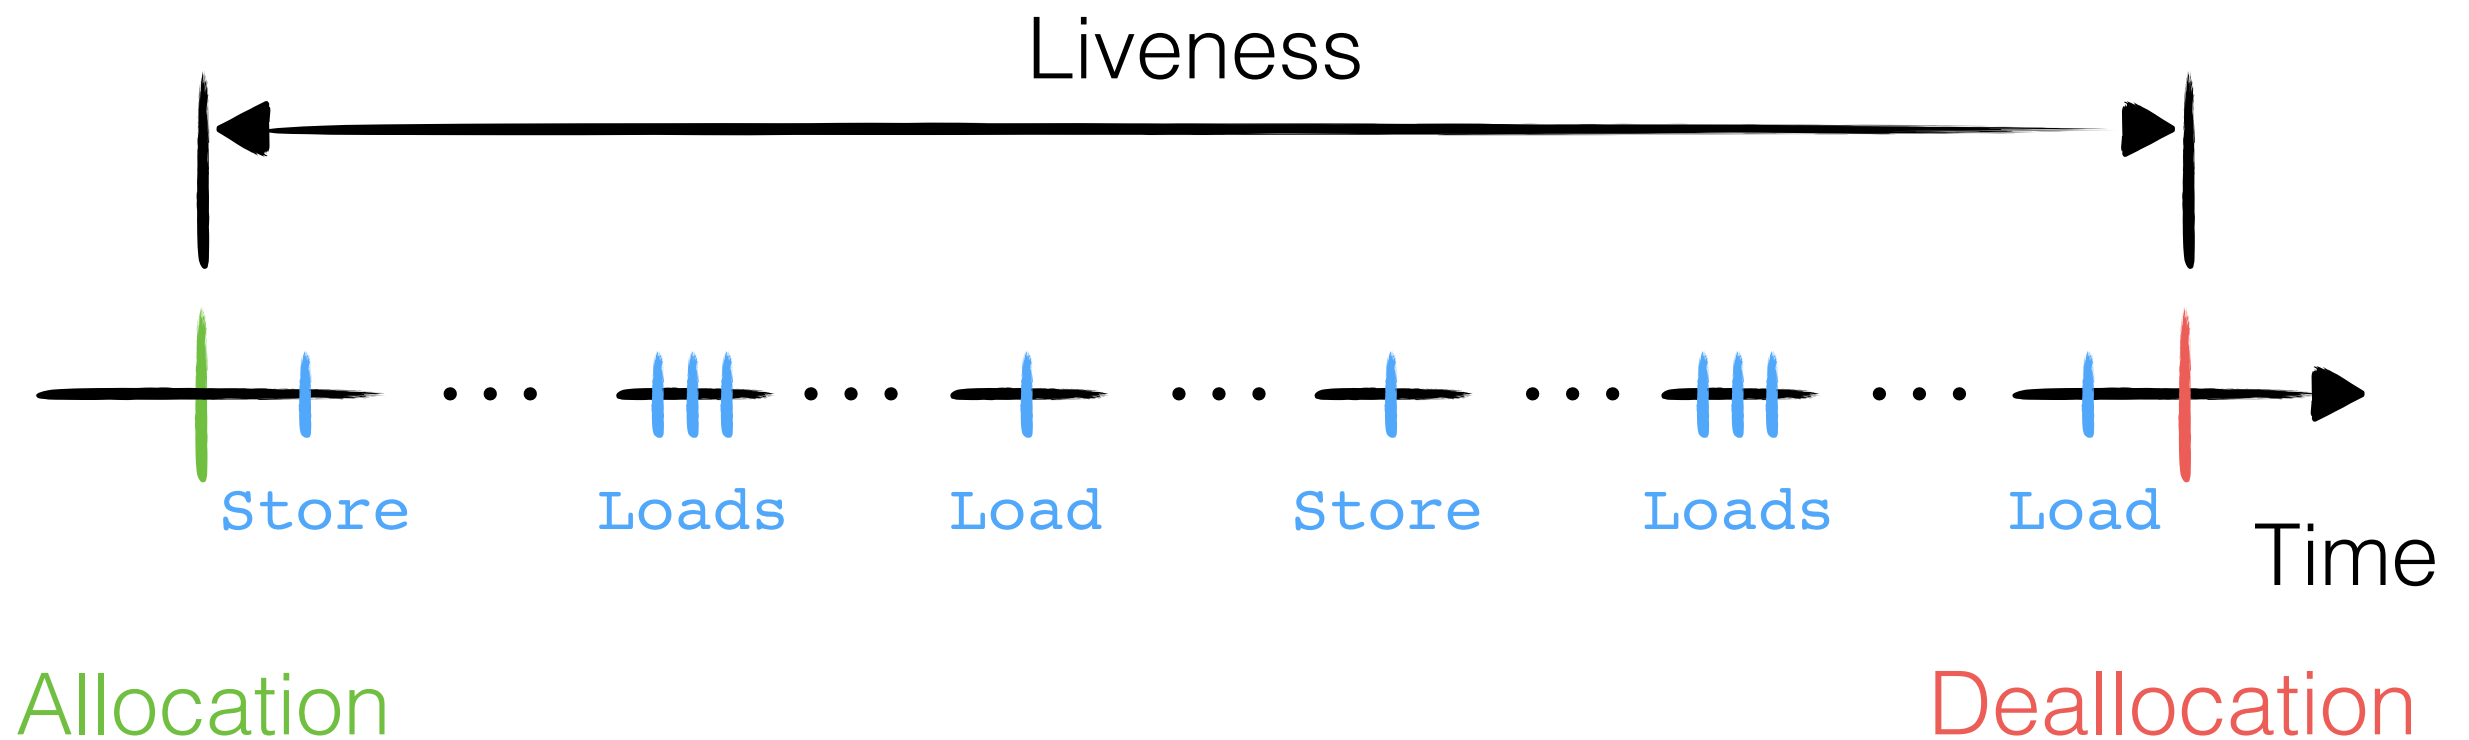
GCC
v4.8.5

Valgrind
Lackey

store &sum
store &x
store &y
load &sum
load &x
store &sum
store &z
load &sum
load &y
store &sum
load &sum
load &z
store &sum

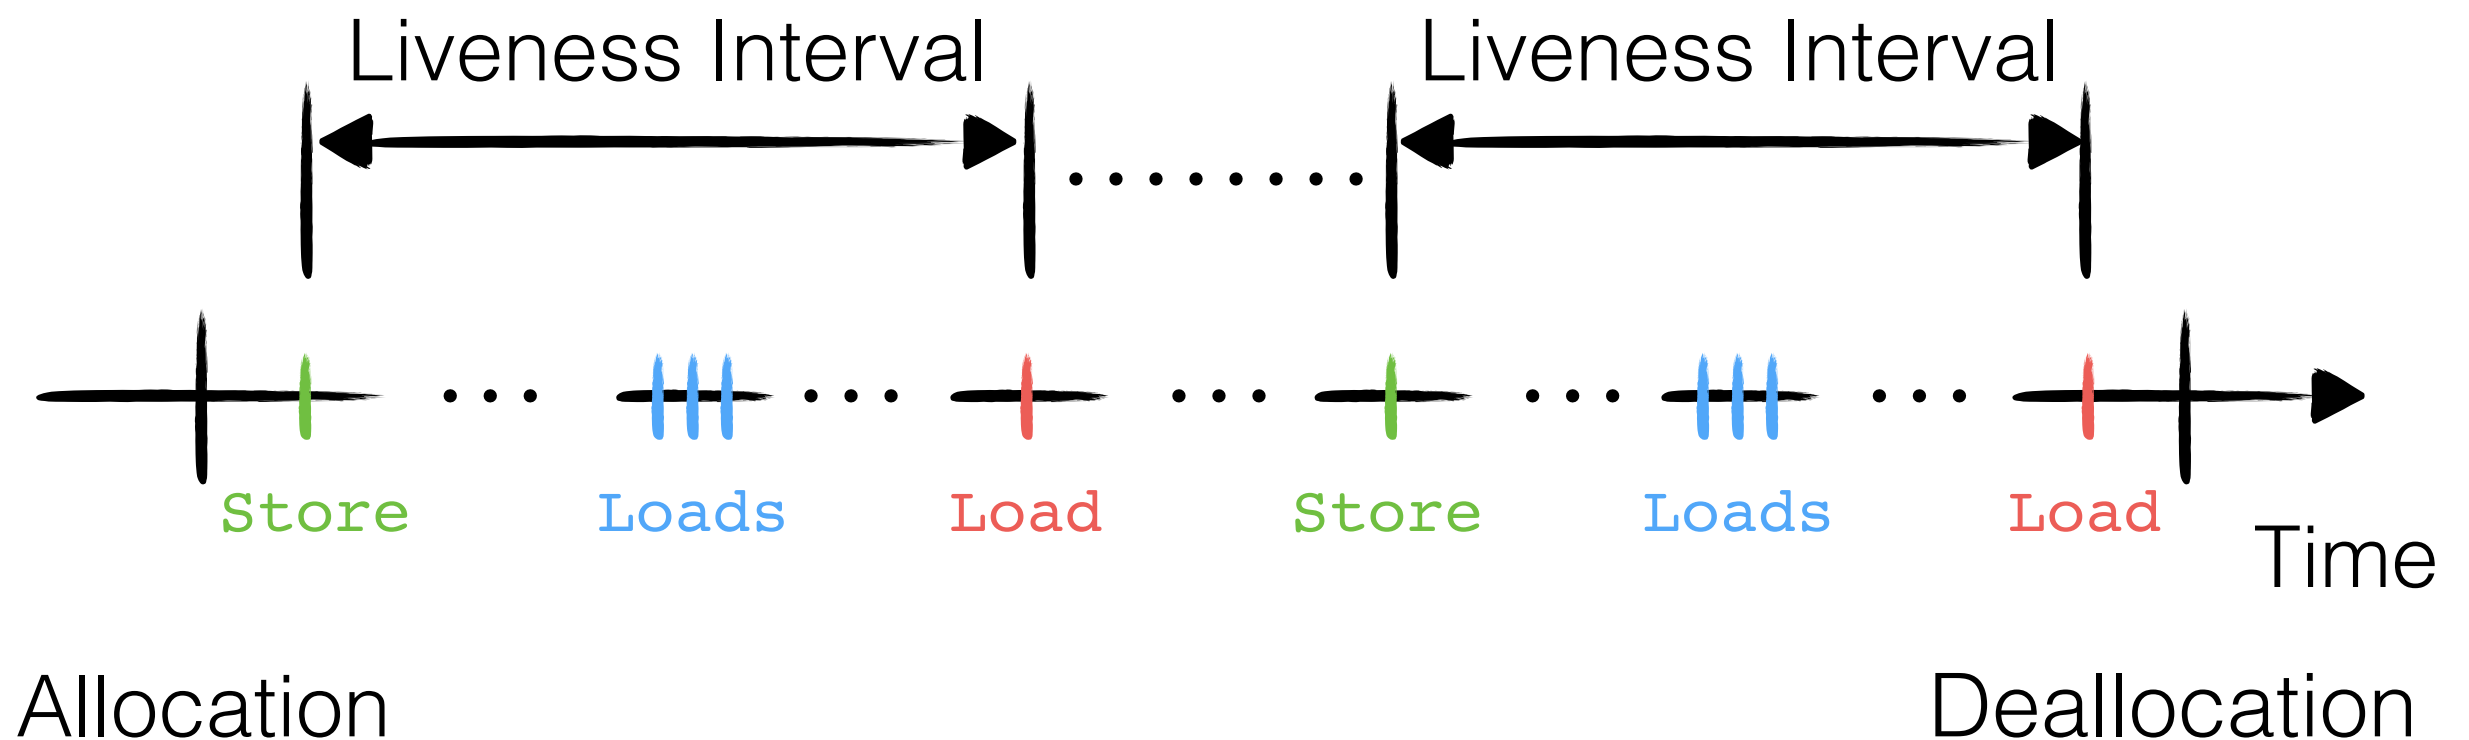
Analysis

Liveness



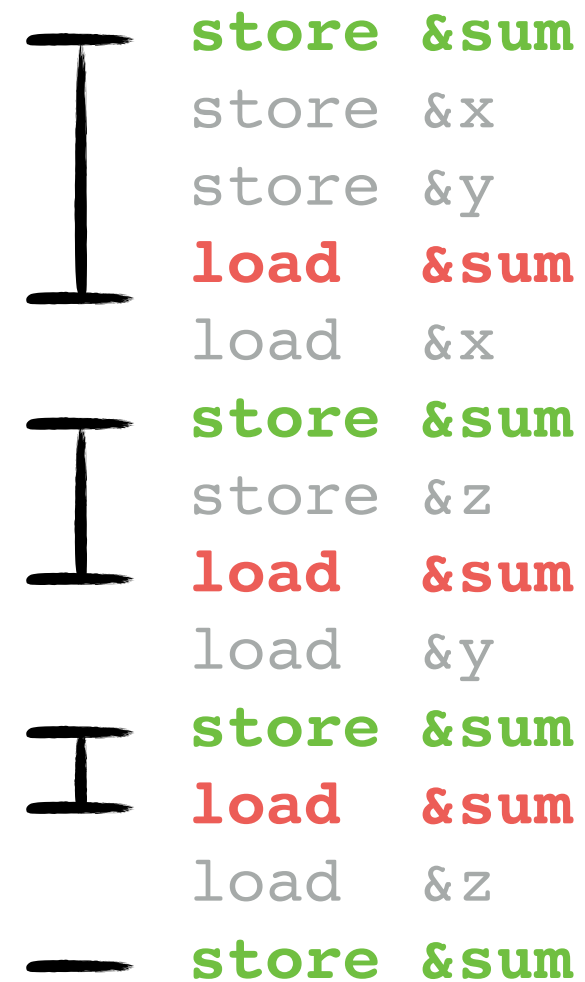
Analysis

Liveness Intervals (LI)



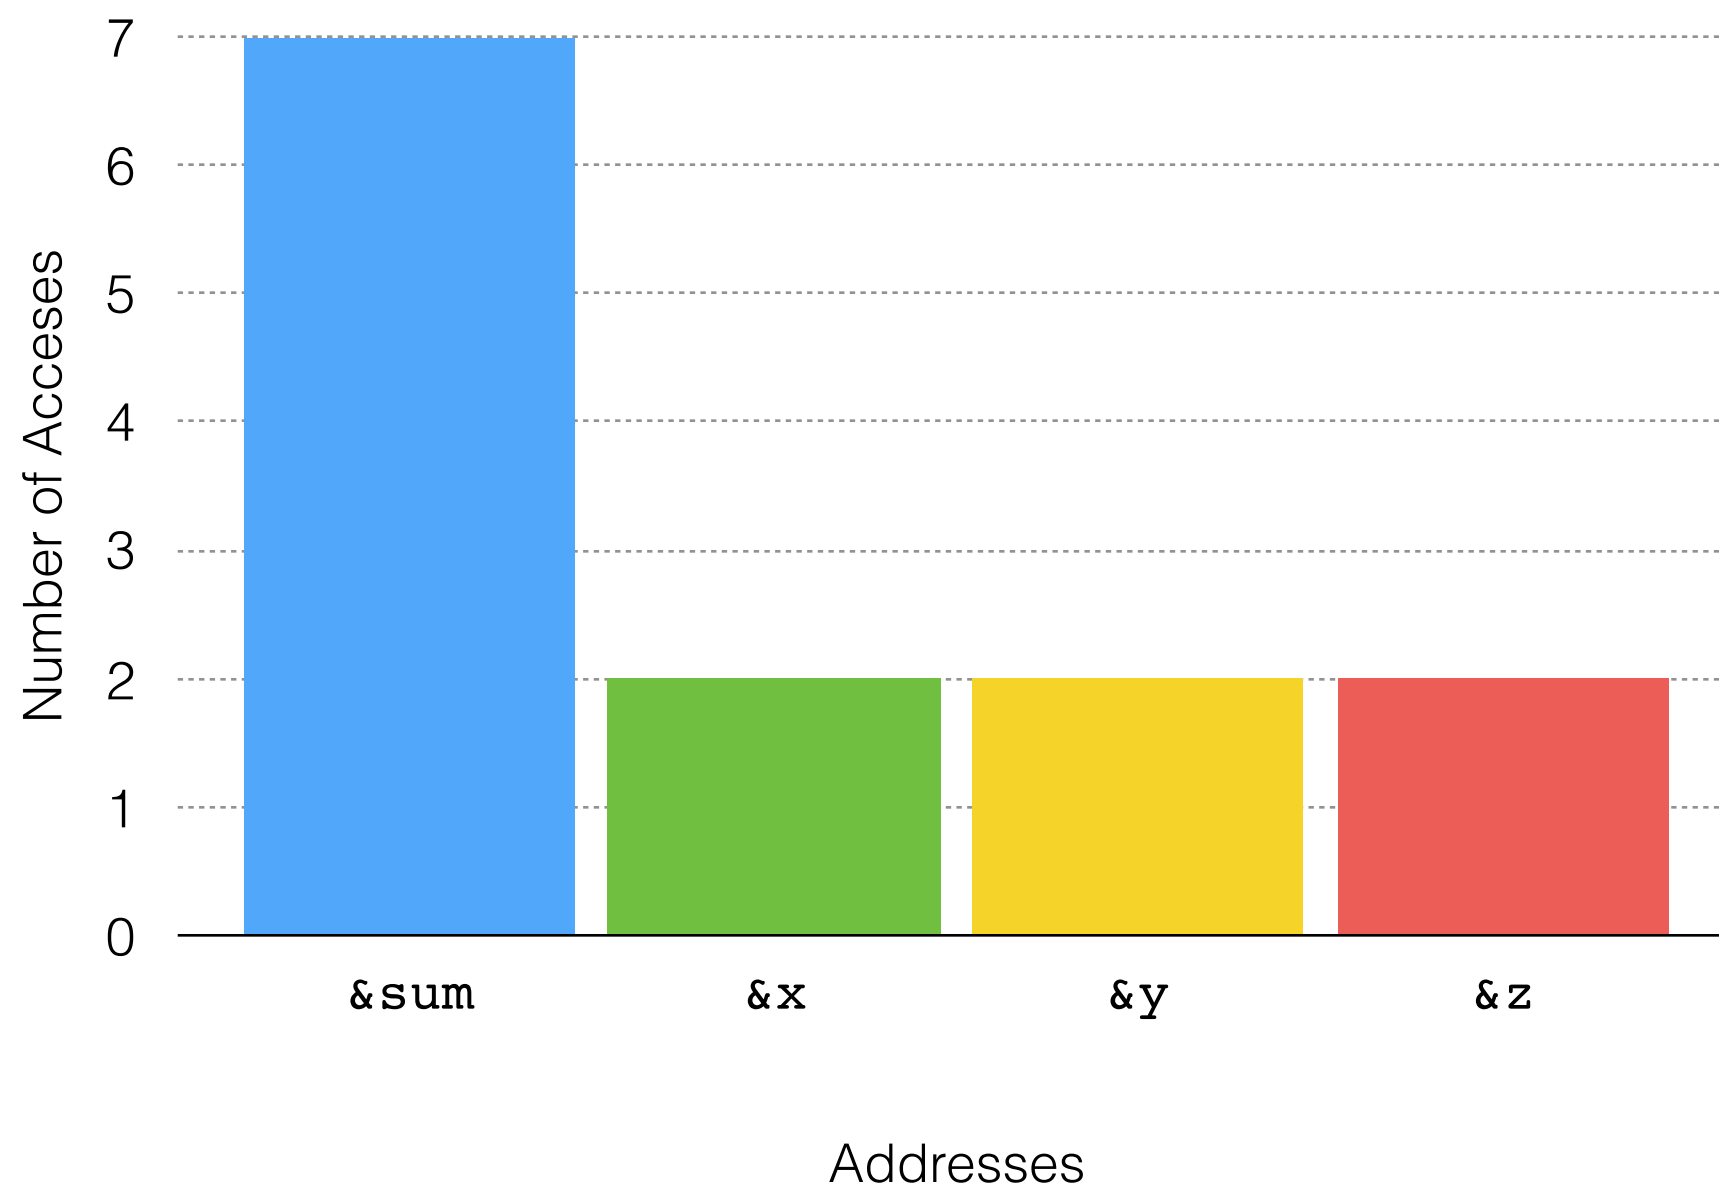
Analysis

Liveness Intervals (LI)



Analysis

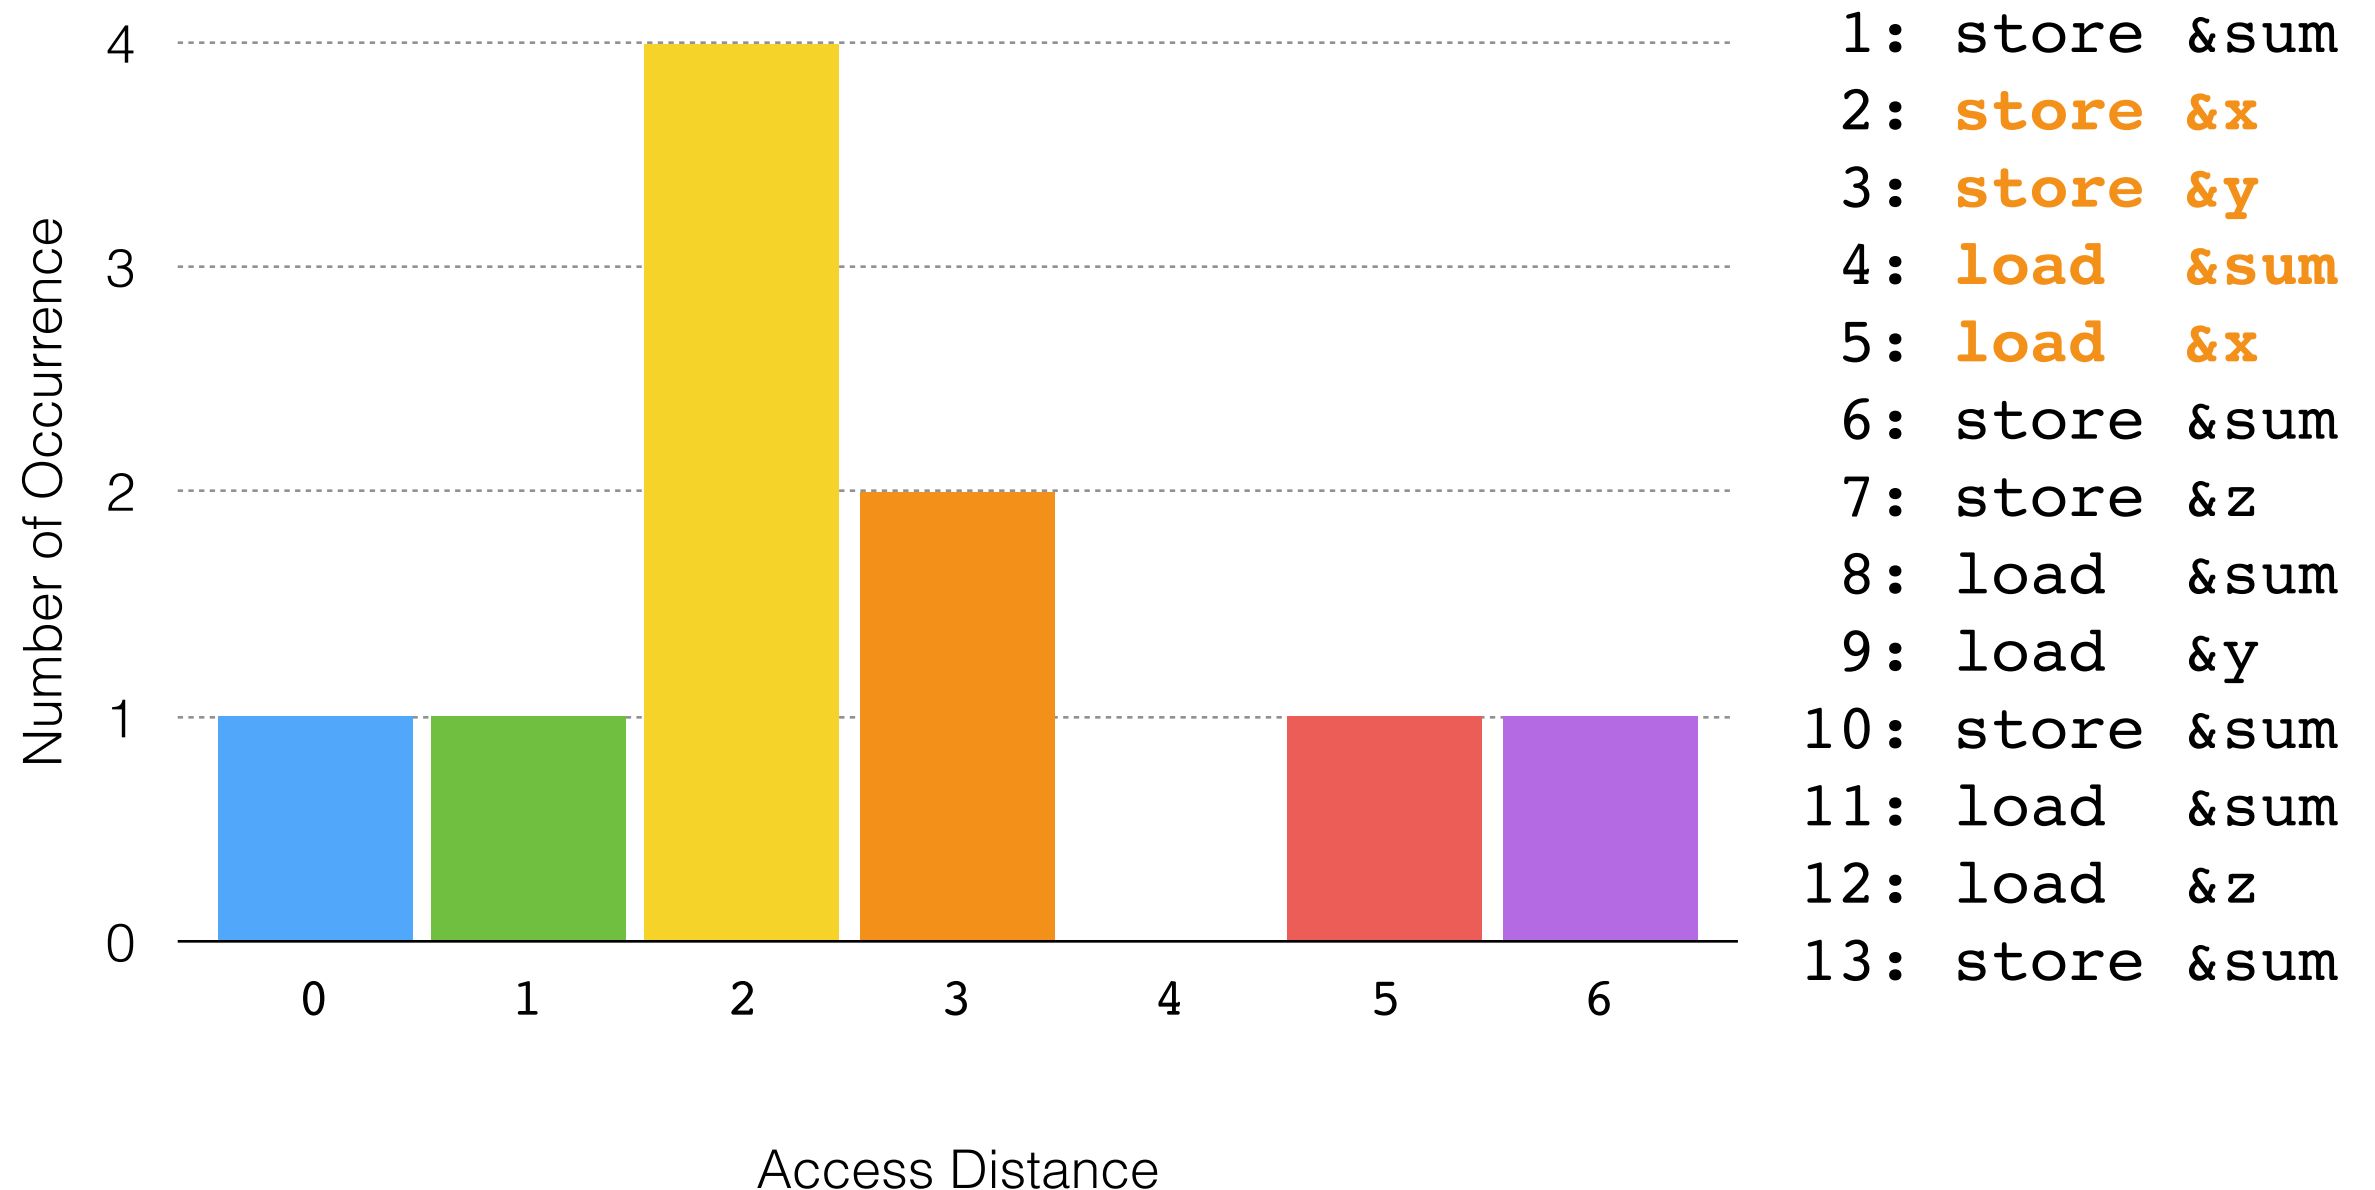
Metric: Accesses



```
1: store &sum
2: store &x
3: store &y
4: load &sum
5: load &x
6: store &sum
7: store &z
8: load &sum
9: load &y
10: store &sum
11: load &sum
12: load &z
13: store &sum
```

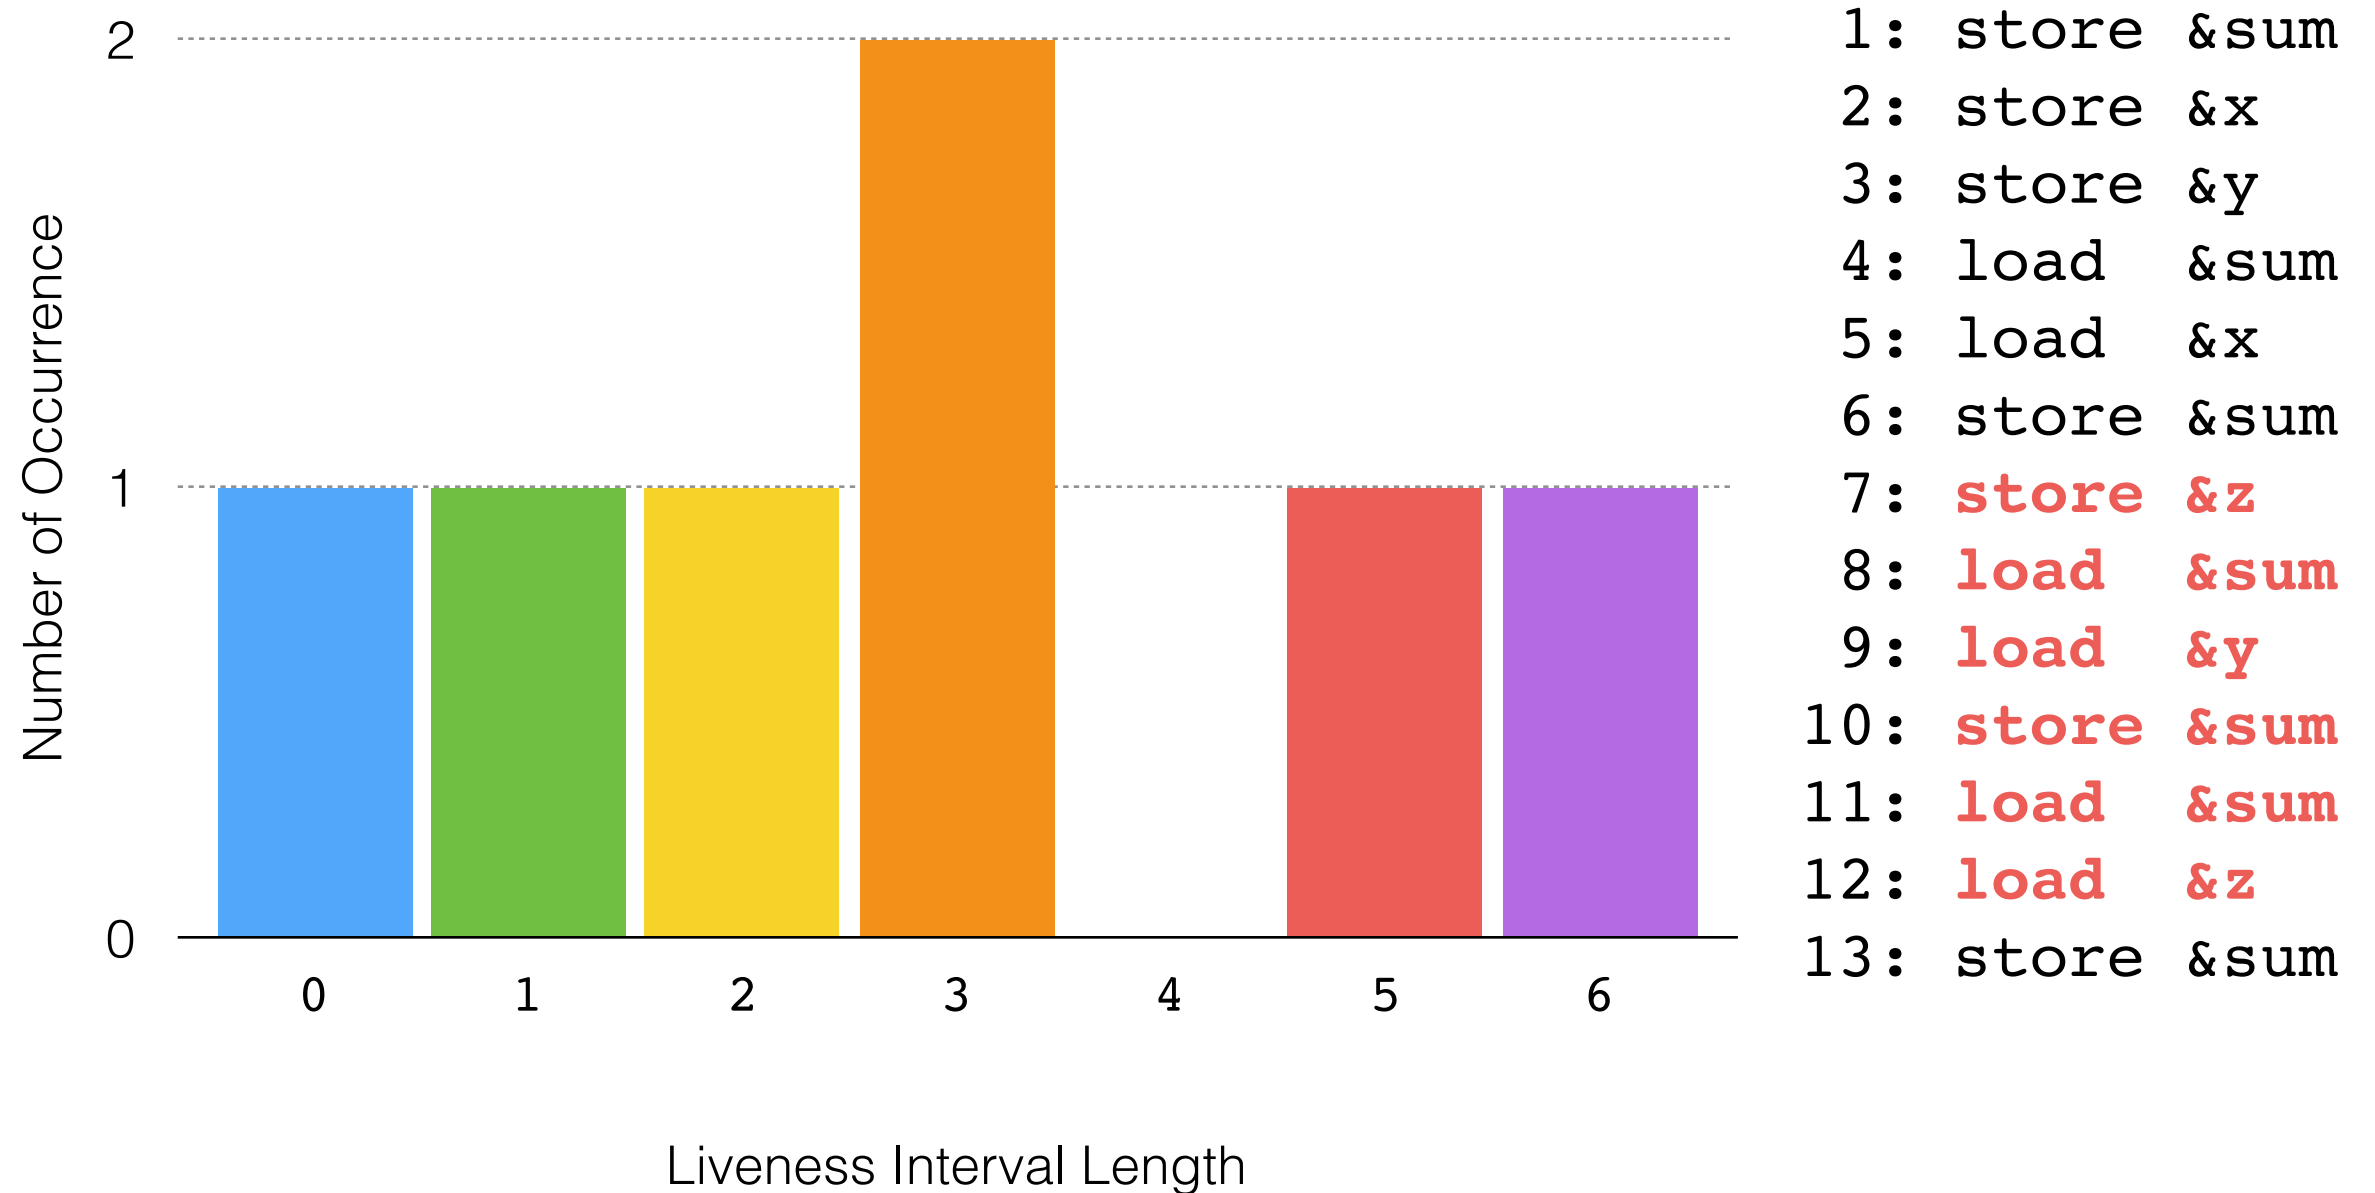
Analysis

Metric: Access Distance



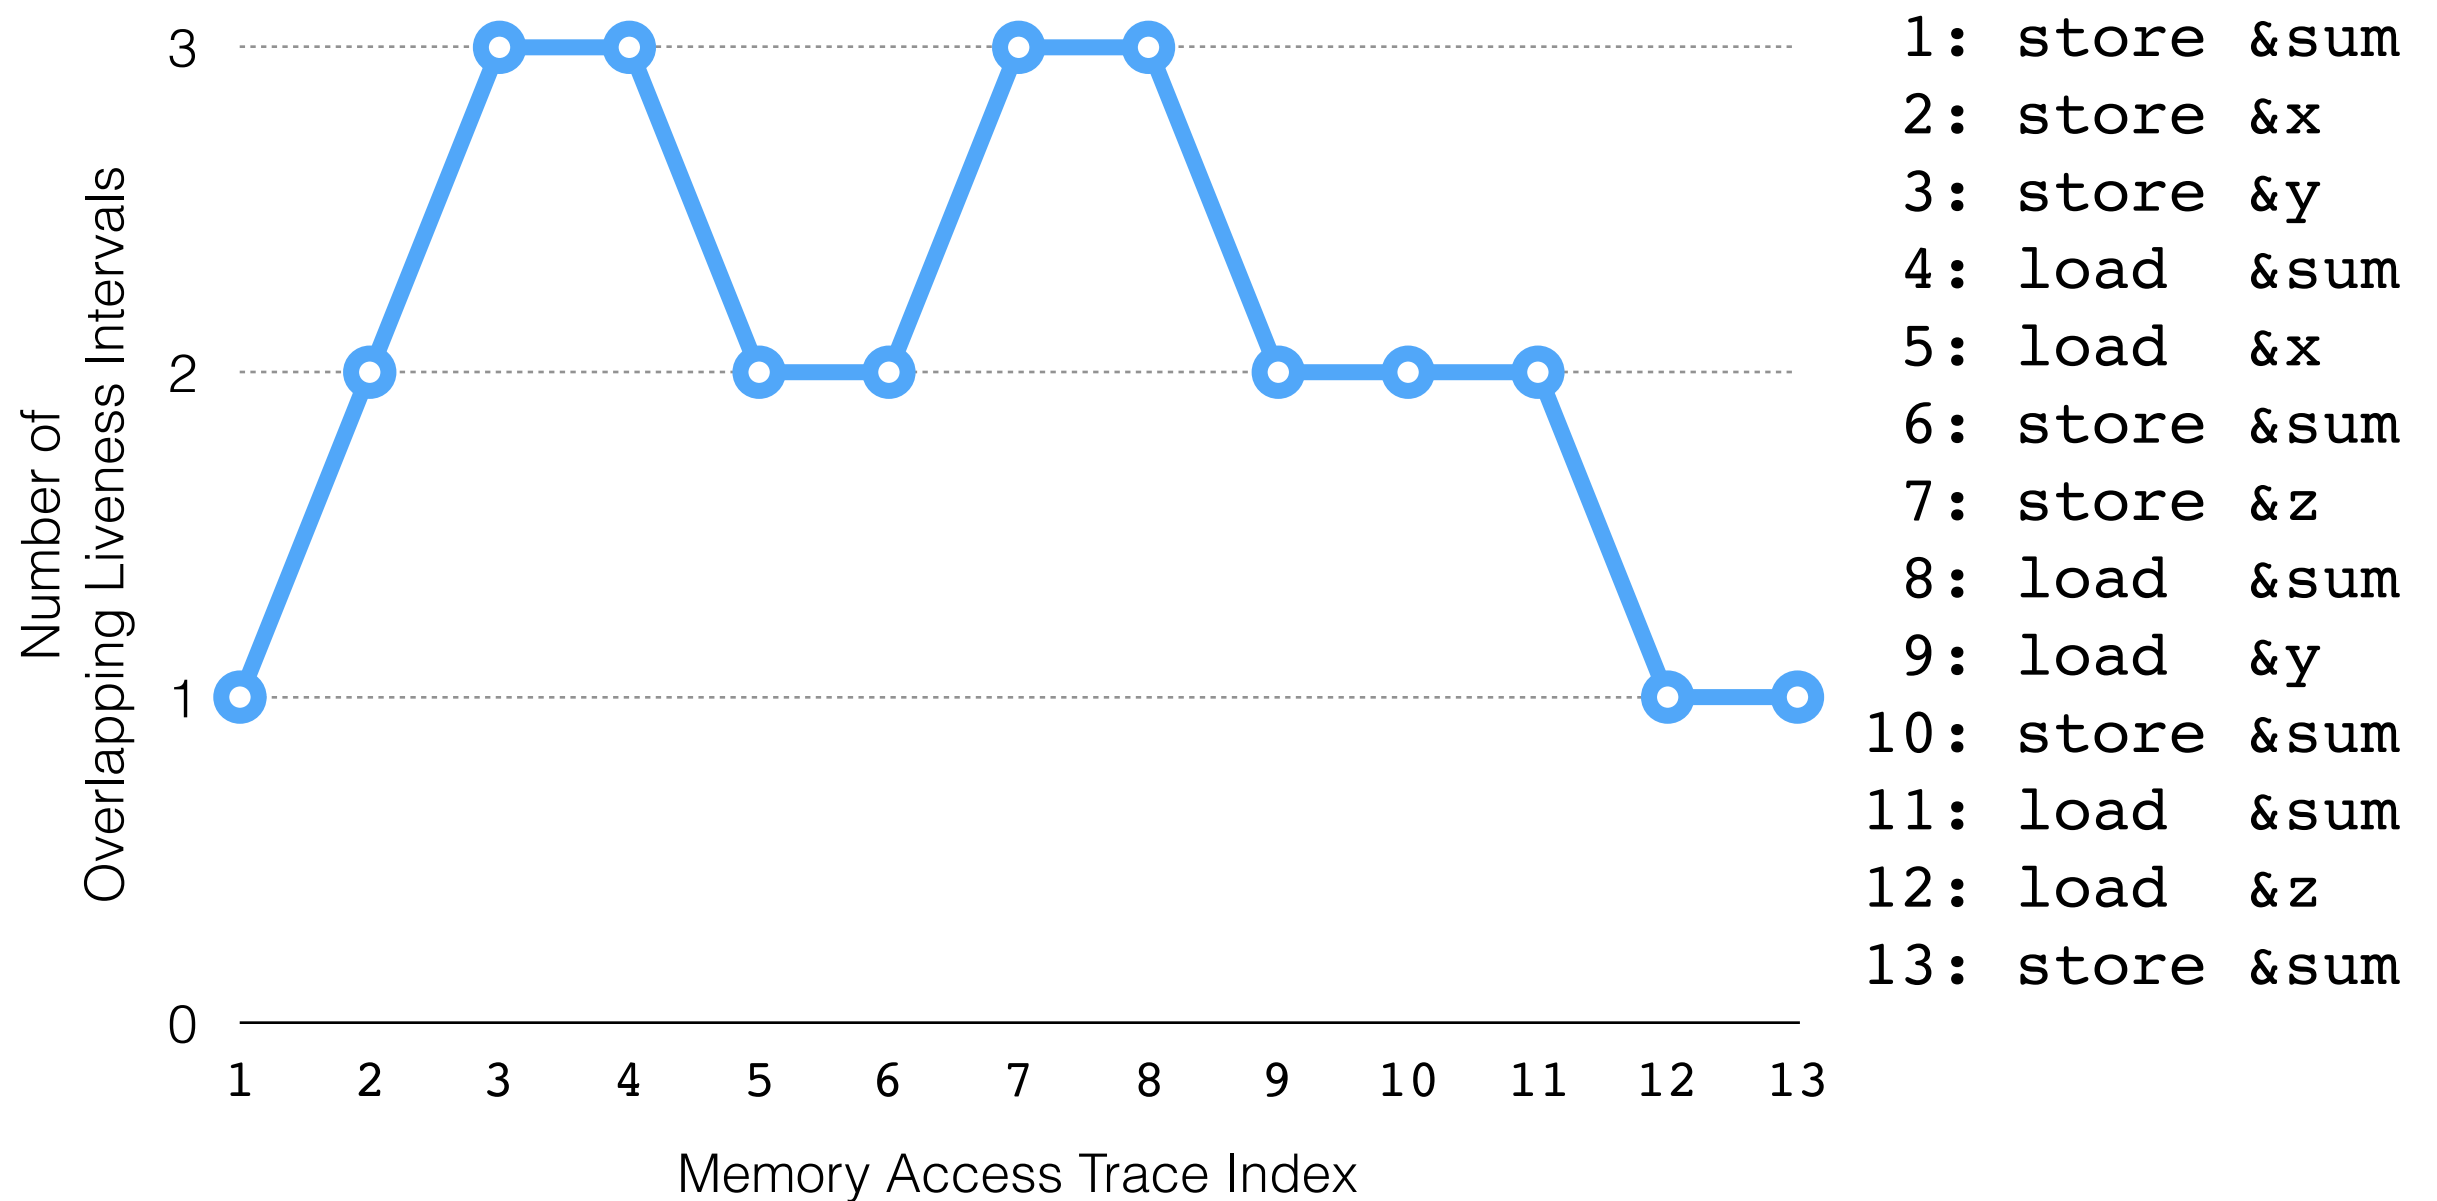
Analysis

Metric: LI Length

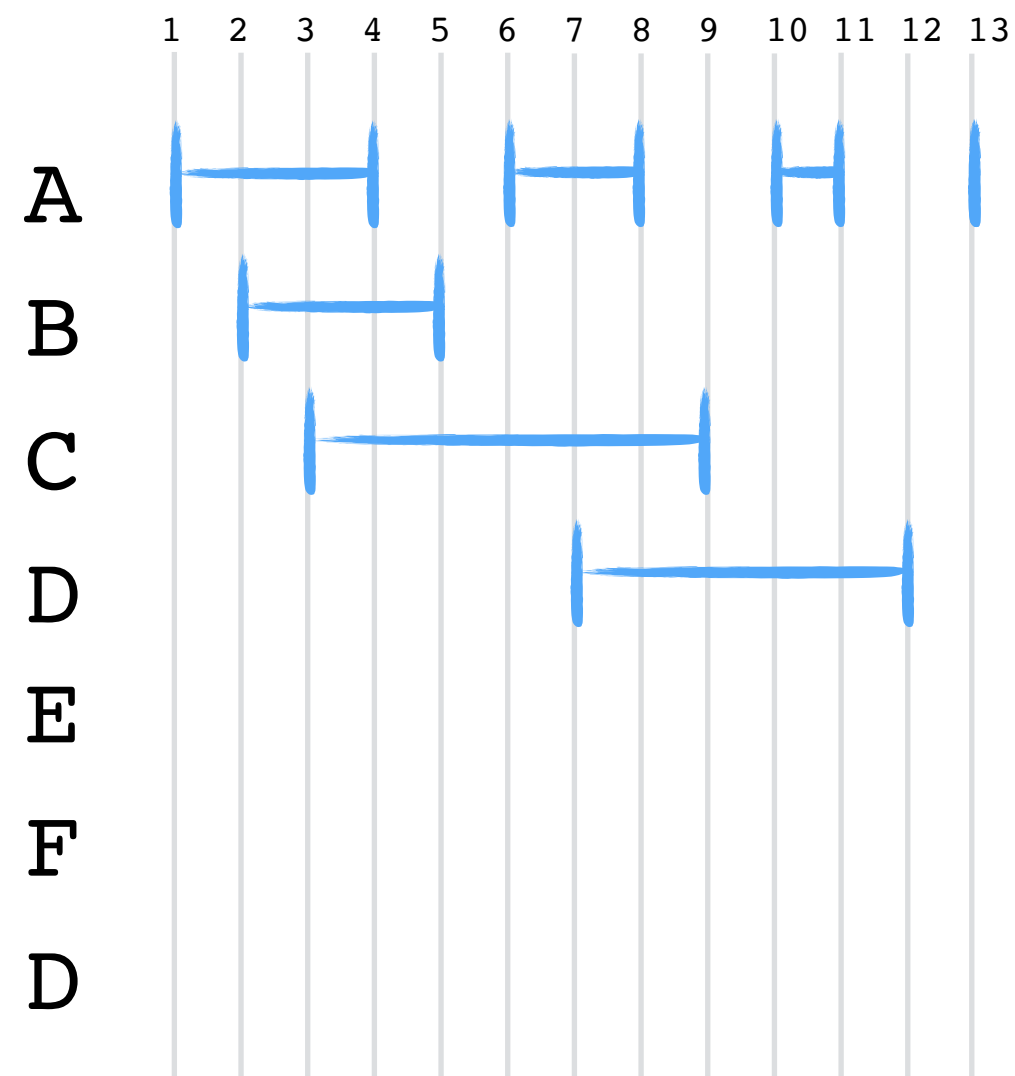


Analysis

Metric: Overlapping LI

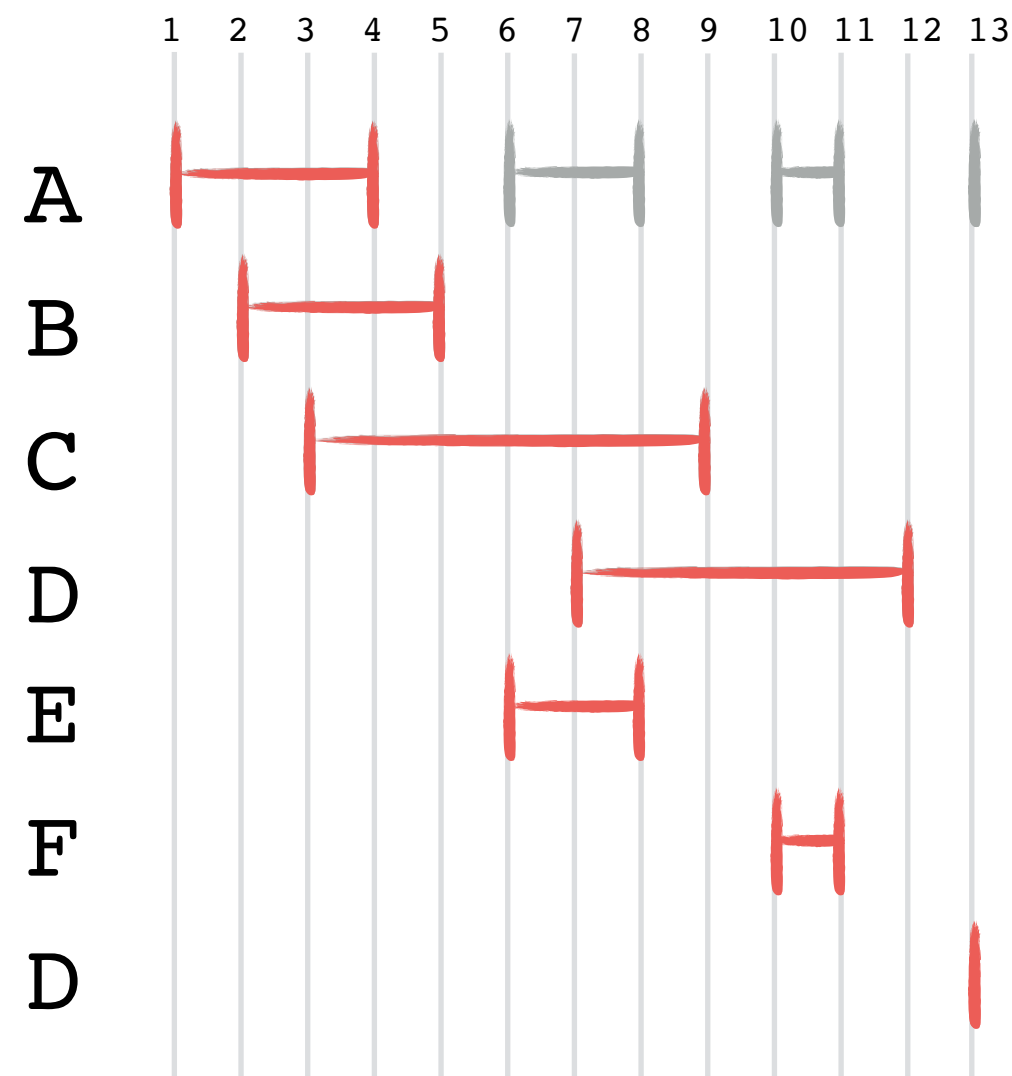


Trace Transformation Identity



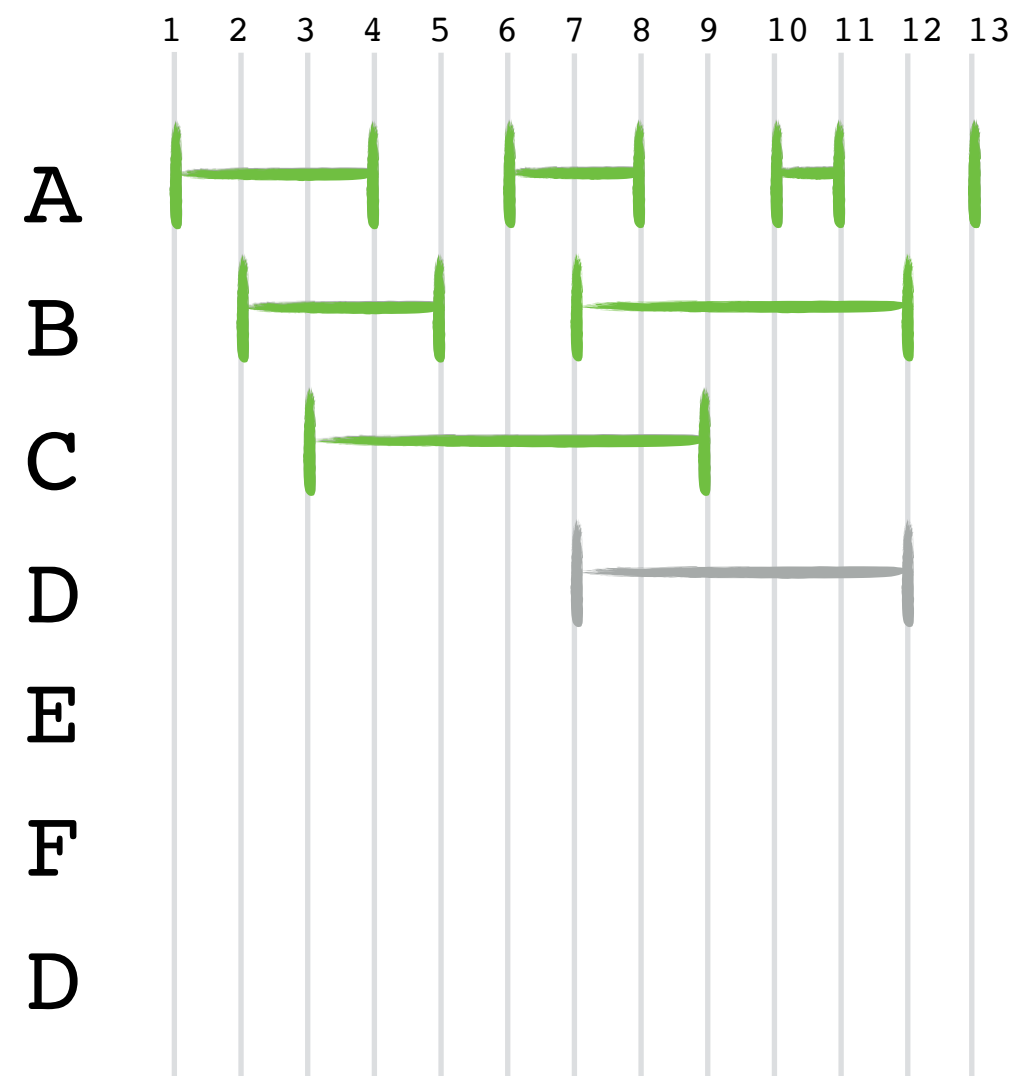
Trace Transformation

Single Assignment

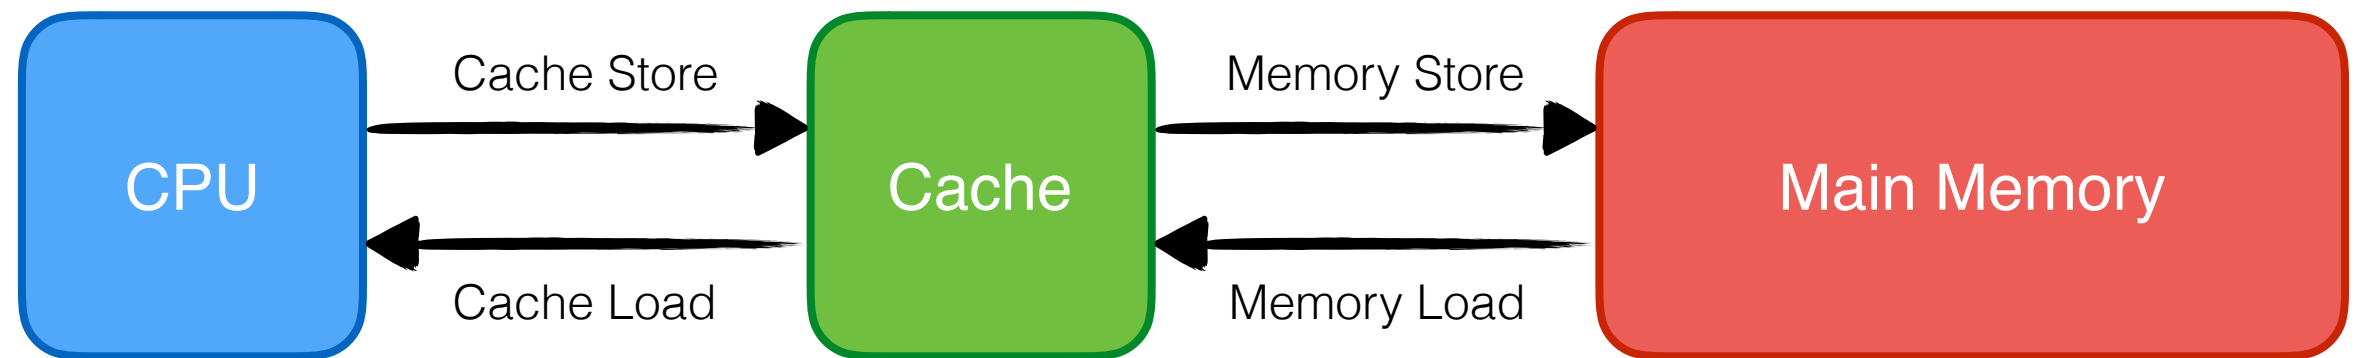


Trace Transformation

Compact



Execution Hardware Model



Execution Performance

$$\text{CPA}(T, C) = \frac{\text{Sum of cycles}(T, C)}{\text{Sum of accesses}(T)}$$

Key

Trace ... T

Cache ... C

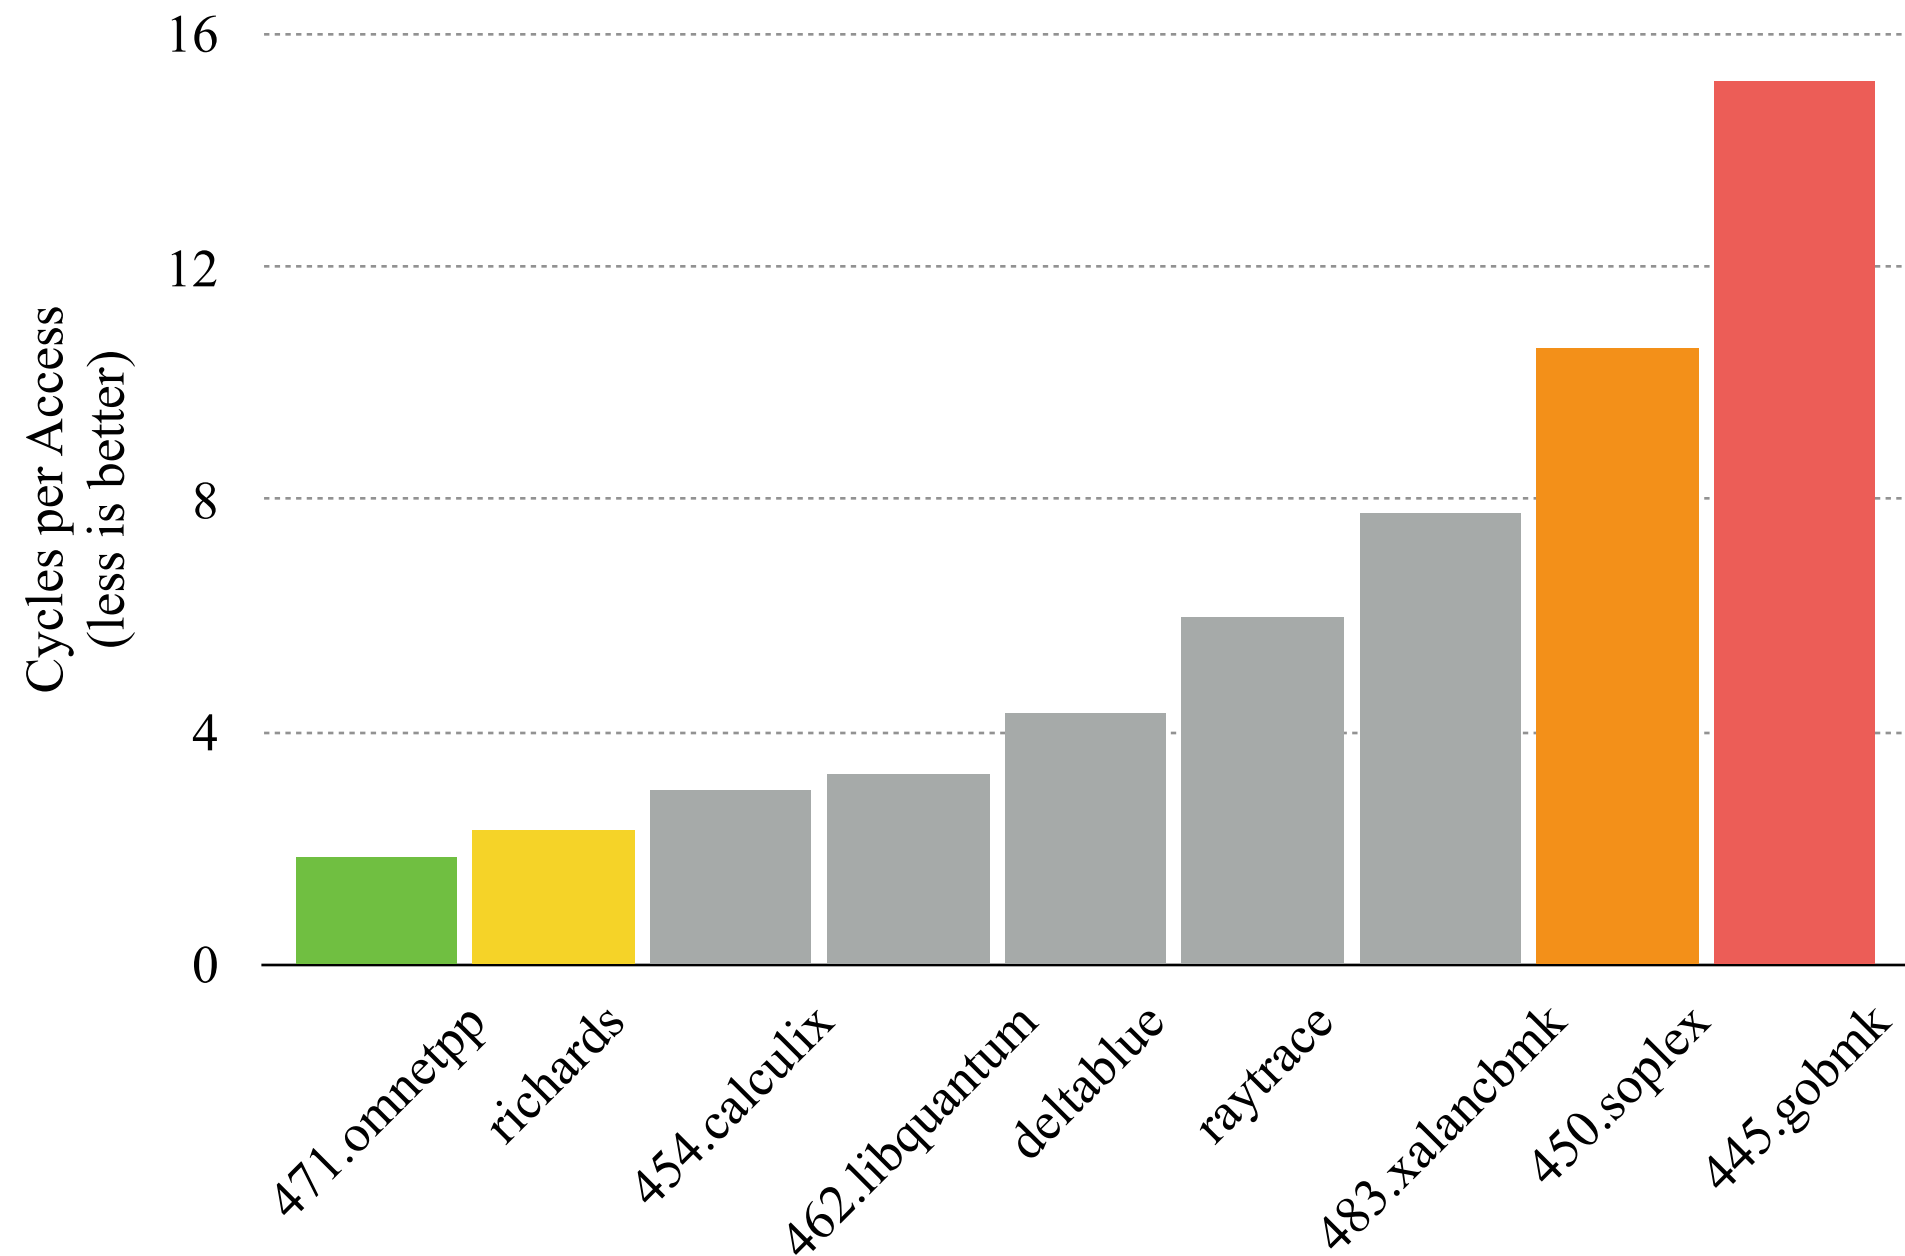
Cycles per Access ... CPA

Costs in Cycles

Cache Load / Store: **1**

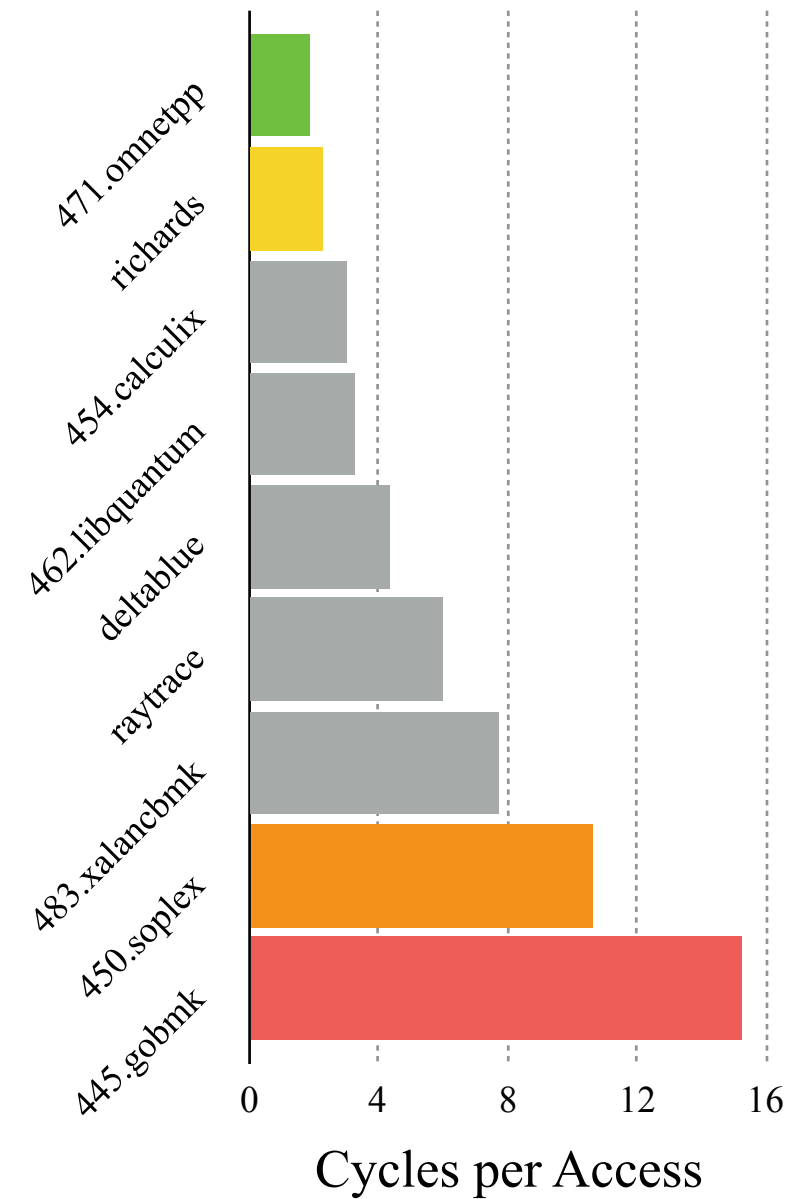
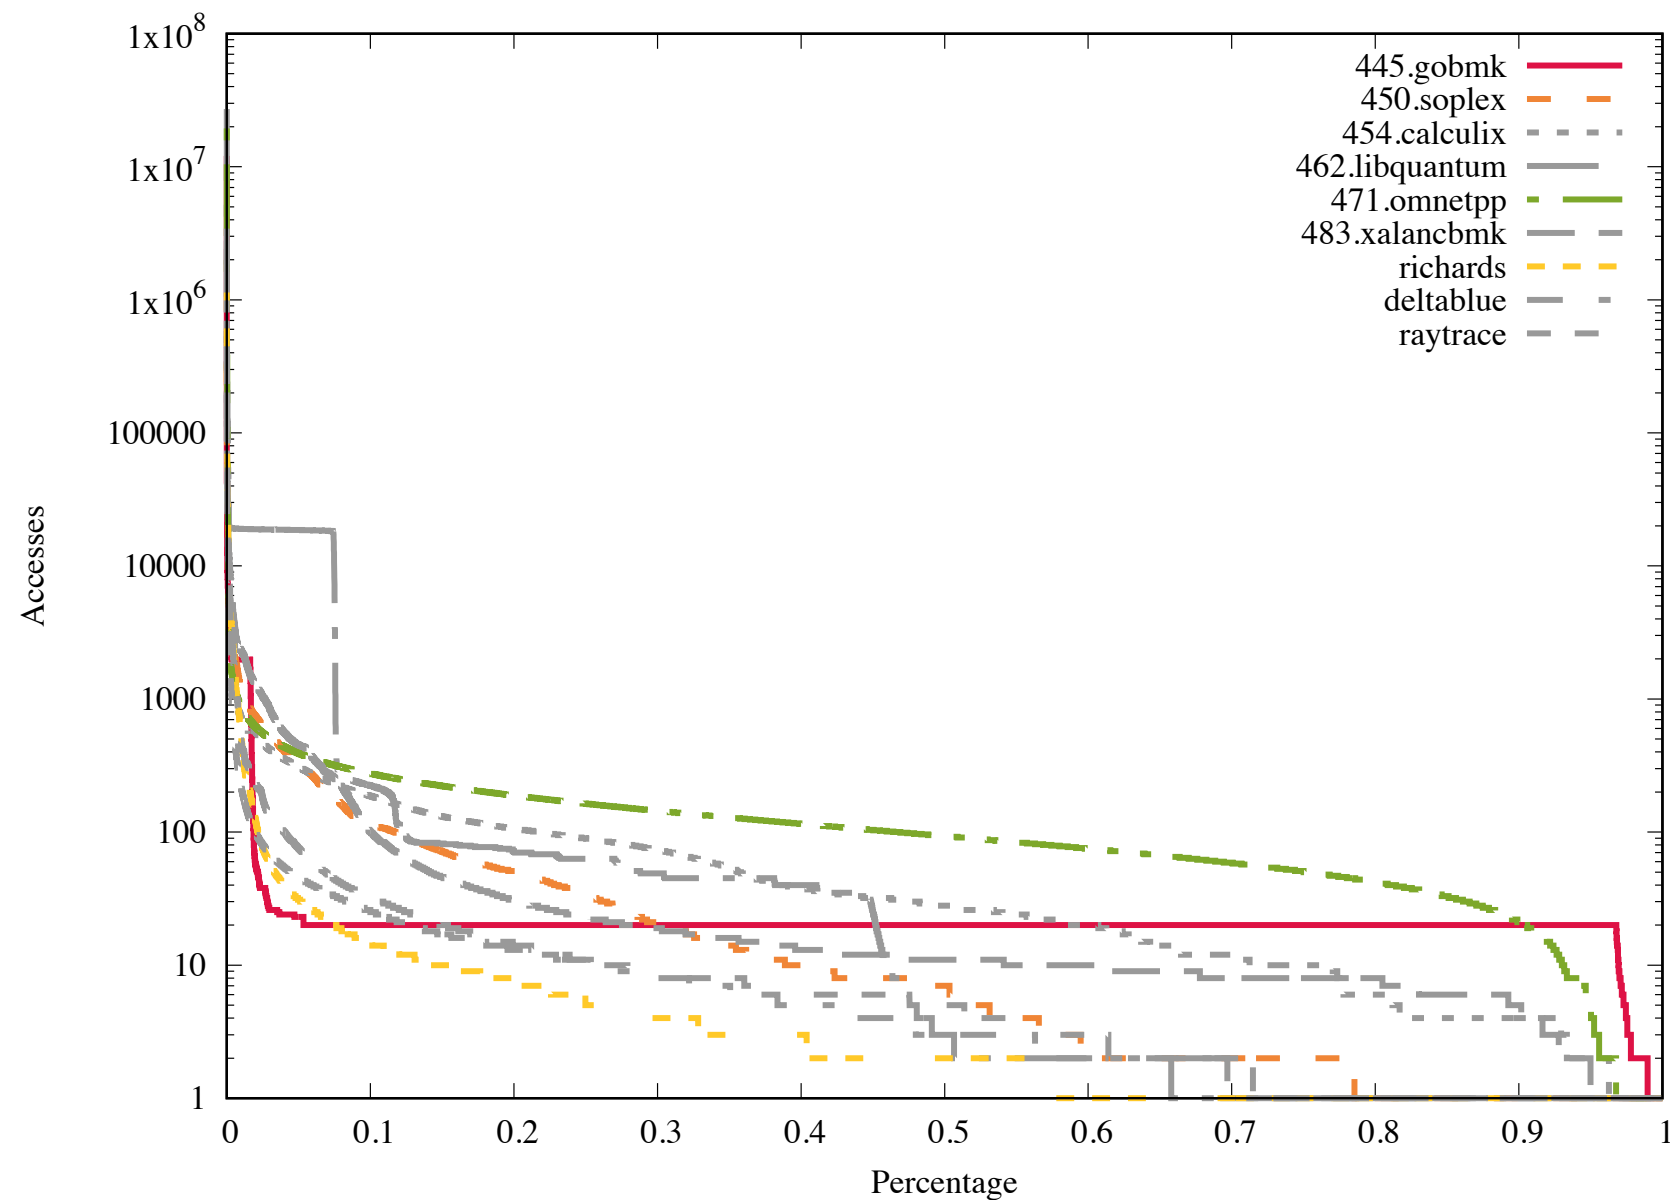
Memory Load / Store: **5**

Experiment Performance



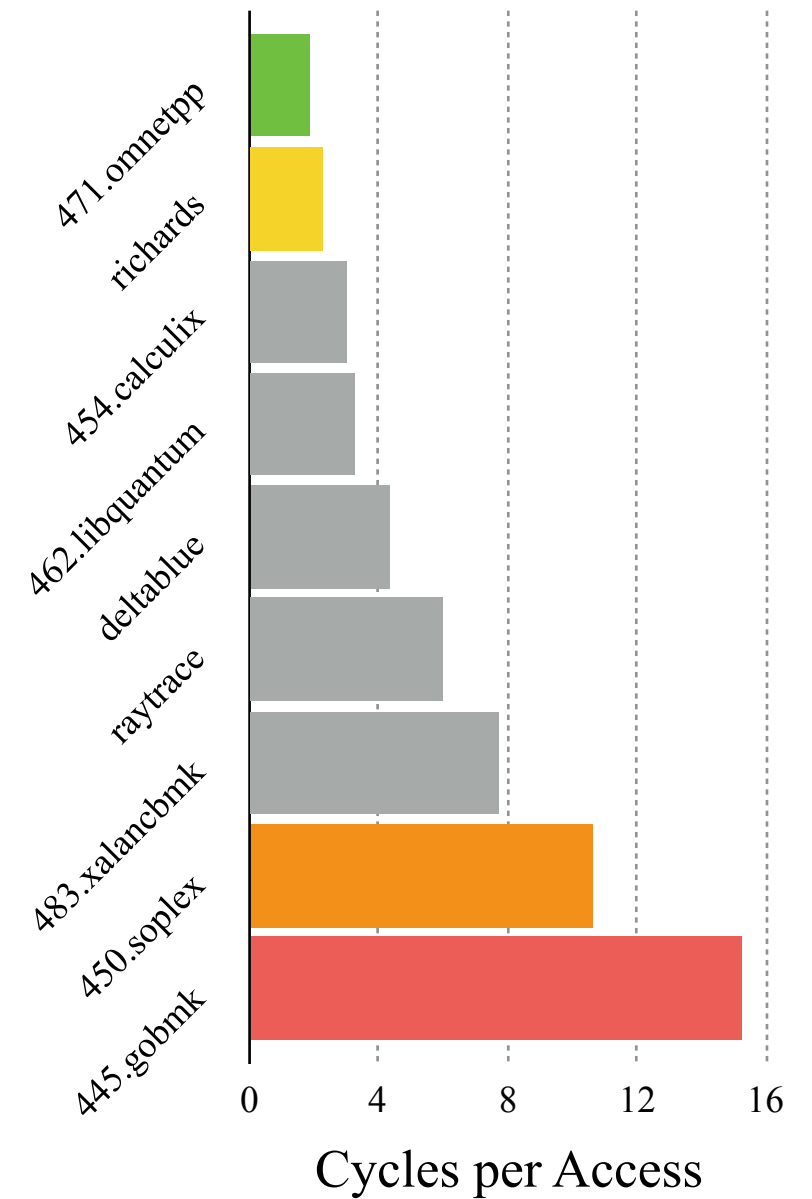
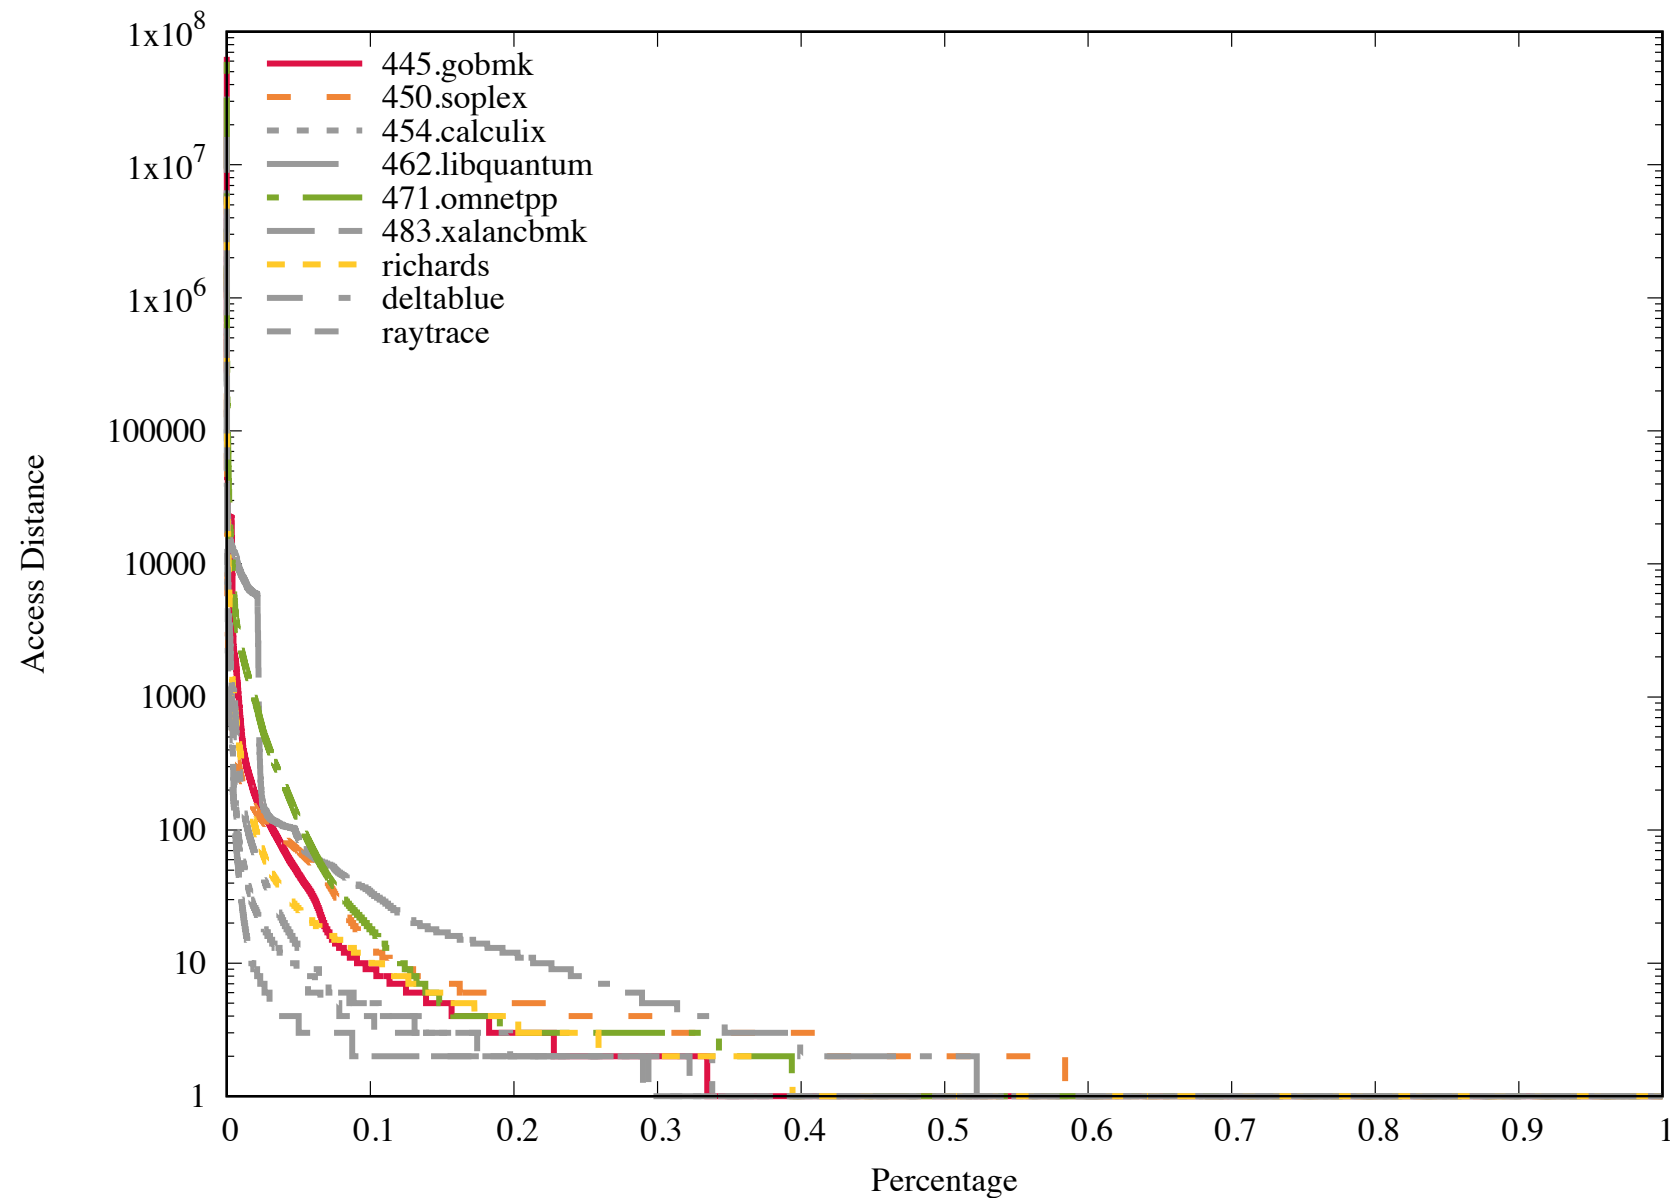
Experiment

Metric: Accesses



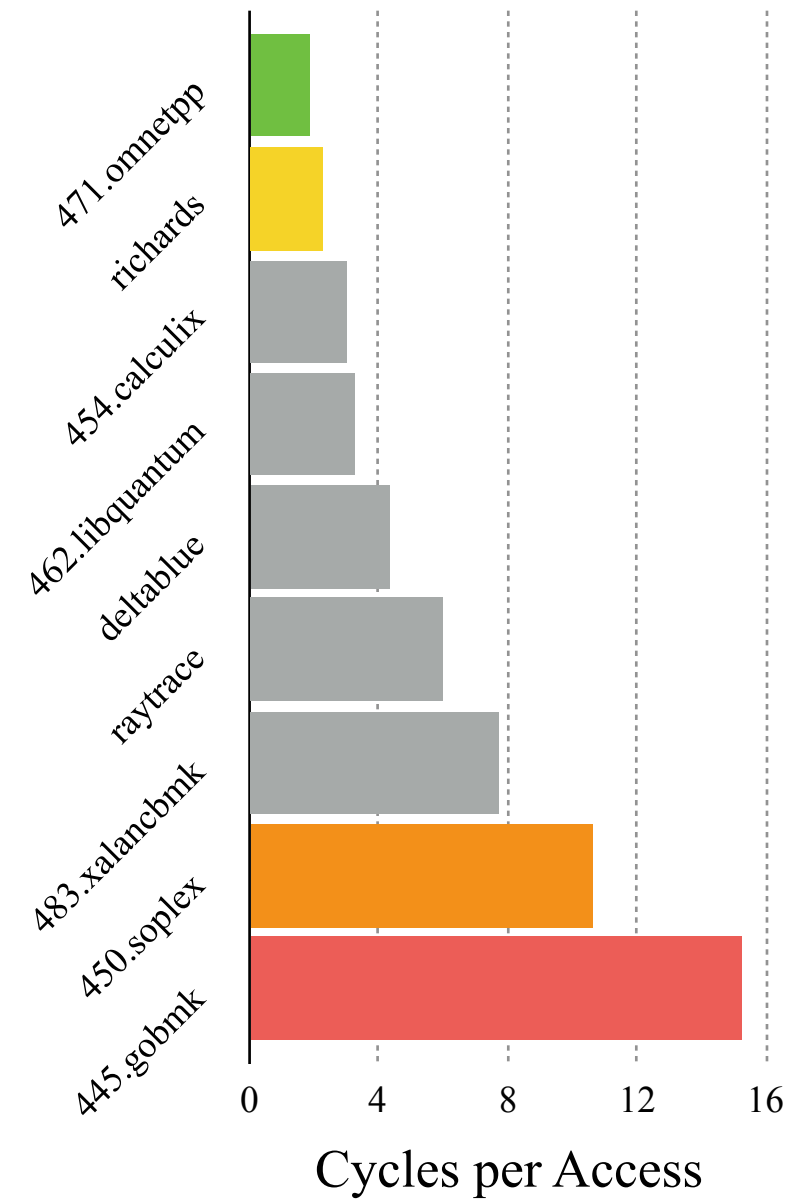
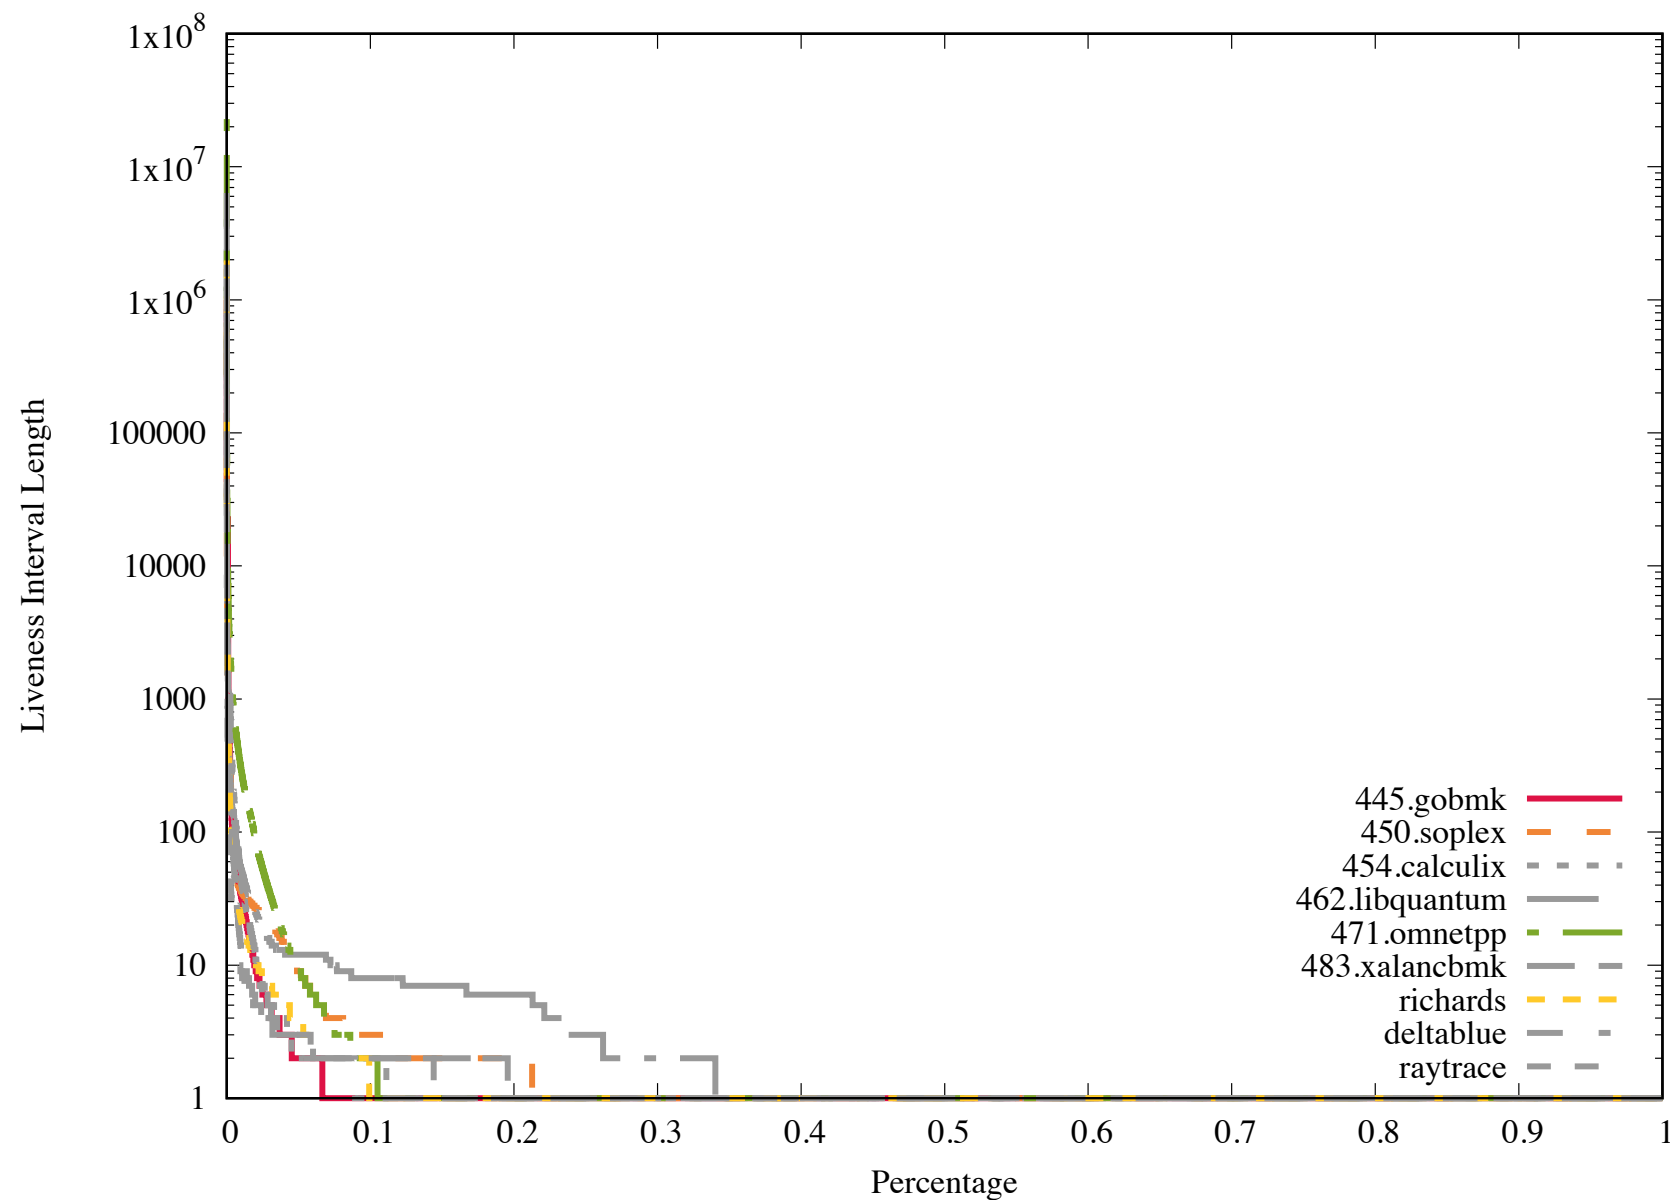
Experiment

Metric: Access Distance



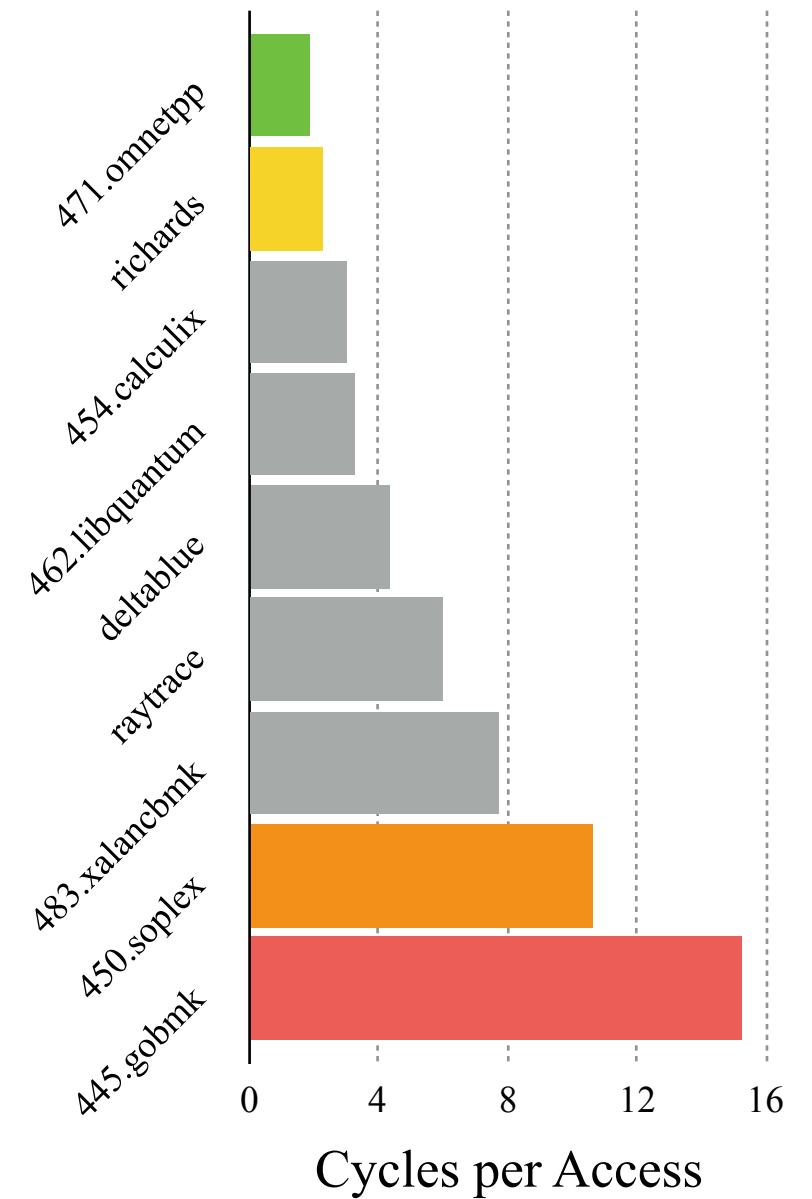
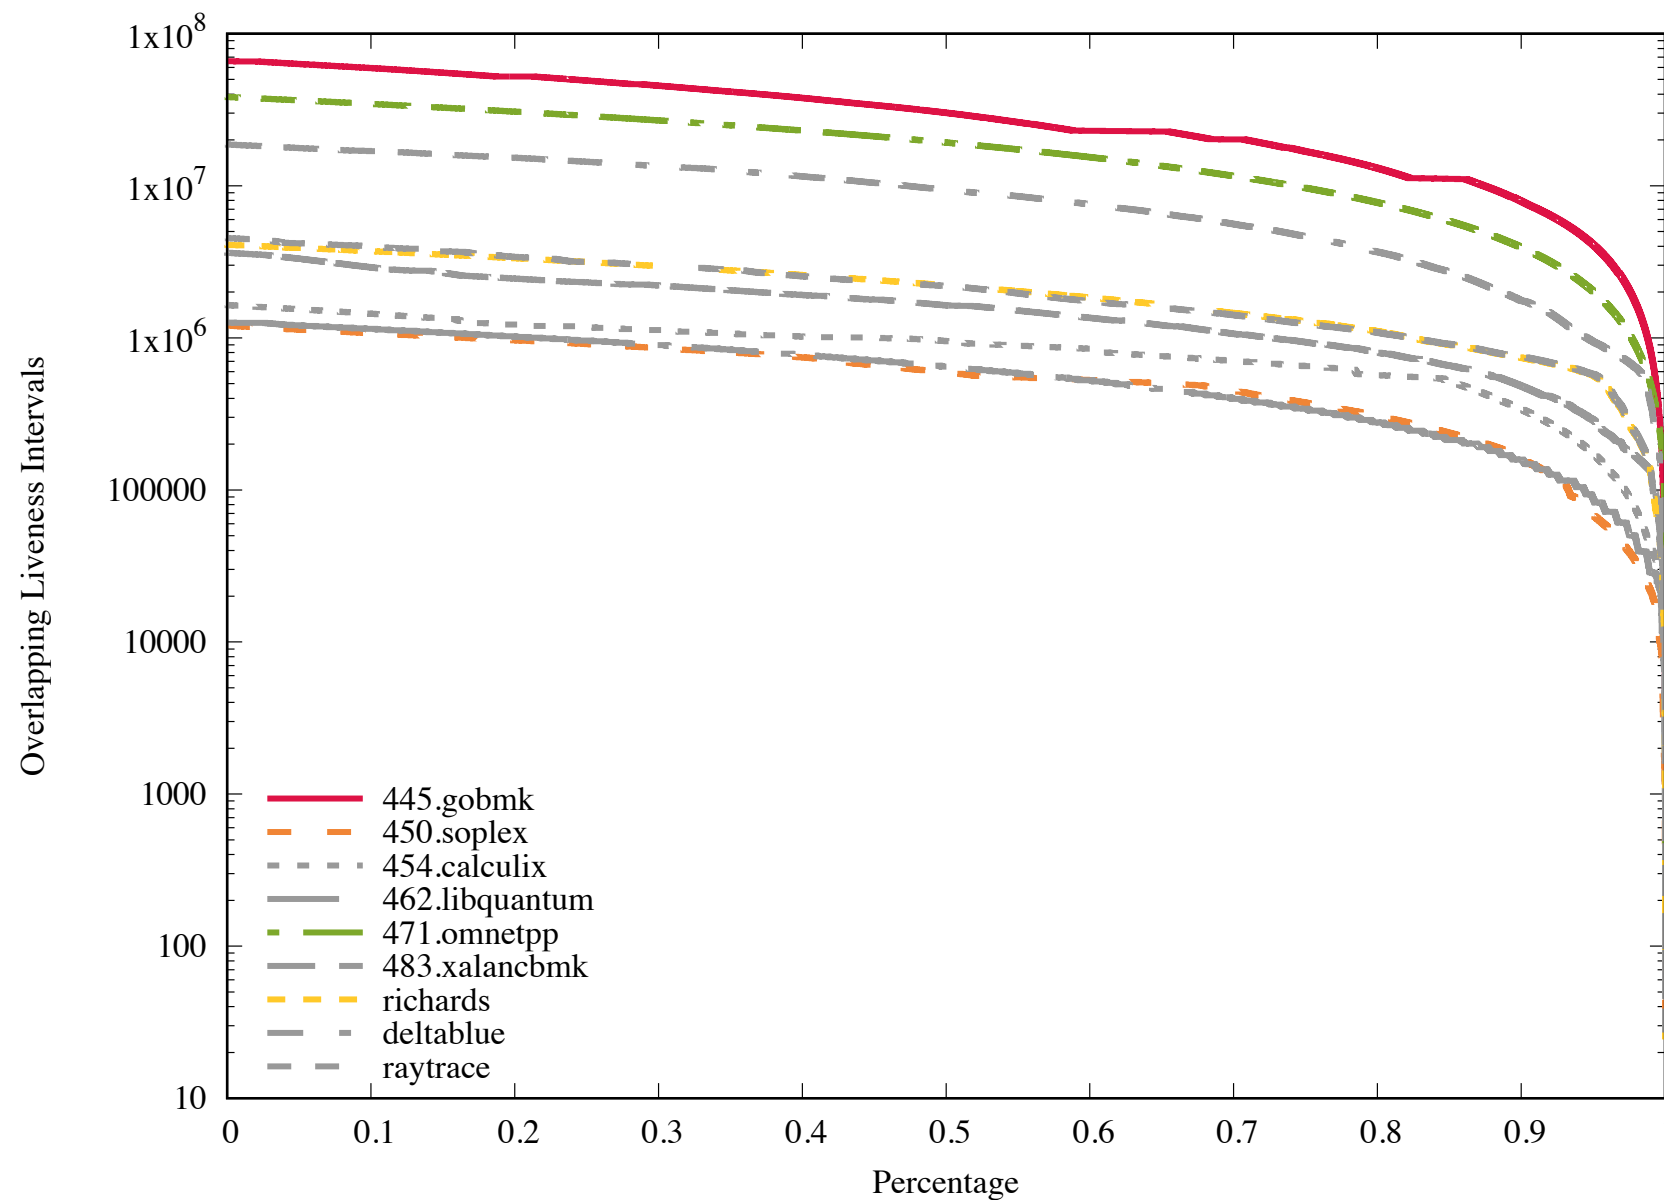
Experiment

Metric: LI Length



Experiment

Metric: Overlapping LI



Conclusion

Metrics

445.gobmk

- Few accesses on the majority of address
- Nearly constant number of accesses
- Many overlapping liveness intervals

471.omnetpp

- Many accesses per address
- Few addresses with significantly more accesses
- Many overlapping liveness intervals

Conclusion

4 metrics

- accesses
- access distance
- liveness interval length
- overlapping liveness

9 benchmarks from 2 benchmark suits

- SPEC 2006
- V8

The defined metrics illustrate tendencies for improvement rather than unique characteristics.

Thank you.

Q&A