



# **Towards cache-optimal address allocation: How fast could your code have run if you had known where to allocate memory?**

MASTER'S THESIS

submitted in partial fulfillment of the requirements for the degree of

**Diplom-Ingenieur**

in

**Computer Sciences**

by

**Mario Preishuber**

Registration Number 01120643

to the Department of Computer Sciences  
at the Faculty of Natural Sciences  
at the Paris Lodron University of Salzburg

Supervisor: Univ.-Prof. Dr. Christoph Kirsch

Salzburg, December 14, 2017

---

Mario Preishuber

Department of Computer Sciences • Faculty of Natural Sciences  
Paris Lodron University of Salzburg  
Jakob-Haringer-Straße 2 • A-5020 Salzburg • +43 (0) 662 8044-6328  
[www.cosy.sbg.ac.at](http://www.cosy.sbg.ac.at)

---

Univ.-Prof. Dr. Christoph Kirsch

---

# Statement of Authentication

I hereby declare that I have written the present thesis independently, without assistance from external parties and without use of other resources than those indicated. The ideas taken directly or indirectly from external sources (including electronic sources) are duly acknowledged in the text. The material, either in full or in part, has not been previously submitted for grading at this or any other academic institution.

Salzburg, December 14, 2017

---

Mario Preishuber



# Acknowledgments

**TODO:** Enter text.  
**TOREVISE**



# Abstract

**TODO:** Enter text.  
**TOREVISE**






# Contents

<b>Statement of Authentication</b>	<b>i</b>
<b>Acknowledgments</b>	<b>iii</b>
<b>Abstract</b>	<b>v</b>
<b>Contents</b>	<b>vii</b>
<b>1 Introduction</b>	<b>1</b>
<b>2 Theoretical Foundations</b>	<b>2</b>
2.1 Hardware Model . . . . .	2
2.1.1 Central Processing Unit . . . . .	2
2.1.2 Main Memory . . . . .	3
2.1.3 Cache . . . . .	3
2.2 Memory Access Trace . . . . .	4
2.3 Liveness . . . . .	6
2.4 Trace Transformation . . . . .	6
2.5 Performance . . . . .	6
2.6 Problem Statement . . . . .	6
<b>3 Experimental Setup</b>	<b>7</b>
3.1 Caches . . . . .	7
3.1.1 Belady Cache . . . . .	7
3.1.2 Belady Cache with Liveness Information . . . . .	7
3.1.3 Least Recently Used Cache . . . . .	7
3.1.4 Least Recently Used Cache with Liveness Information . . . . .	7
3.2 Allocators . . . . .	7
3.2.1 Original Allocator . . . . .	7
3.2.2 Single Assignment Allocator . . . . .	7
3.2.3 Compacting Allocator . . . . .	7
3.2.3.1 Compacting Allocator with Stack Semantic Free List	7
3.2.3.2 Compacting Allocator with Queue Semantic Free List	7
3.2.3.3 Compacting Allocator with Set Semantic Free List .	7

3.3	Benchmarks . . . . .	7
3.3.1	SPEC 2006 Benchmarks . . . . .	7
3.3.1.1	445.gobmk . . . . .	8
3.3.1.2	450.soplex . . . . .	8
3.3.1.3	454.calculix . . . . .	8
3.3.1.4	462.libquantum . . . . .	8
3.3.1.5	471.omnetpp . . . . .	9
3.3.1.6	483.xalancbmk . . . . .	9
3.3.2	Google JavaScript Engine (V8) Benchmarks . . . . .	10
3.3.2.1	Richards . . . . .	10
3.3.2.2	Raytrace . . . . .	10
3.3.2.3	Deltablue . . . . .	11
3.4	Metrics . . . . .	11
3.4.1	Address Access . . . . .	11
3.4.2	Access Distance . . . . .	11
3.4.3	Live Addresses . . . . .	11
3.4.4	Liveness Length . . . . .	11
3.5	Trace generation (Valgrind / Cachegrind) . . . . .	11
<b>4</b>	<b>Experiments</b>	<b>12</b>
4.1	8 Byte Cache Line Size . . . . .	12
4.2	64 Byte Cache Line Size . . . . .	12
<b>5</b>	<b>Related Work</b>	<b>13</b>
<b>6</b>	<b>Conclusion</b>	<b>14</b>
<b>7</b>	<b>Future Work</b>	<b>15</b>
<b>8</b>	<b>Acronyms</b>	<b>16</b>
	<b>Appendix</b>	<b>17</b>
	<b>List of Figures</b>	<b>18</b>
	<b>List of Tables</b>	<b>19</b>
	<b>List of Algorithms</b>	<b>20</b>
	<b>Acronyms</b>	<b>21</b>
	<b>Bibliography</b>	<b>22</b>

CHAPTER 1



# Introduction

# Theoretical Foundations

## 2.1 Hardware Model

This section deals with the hardware model applied. The used model is reduced to the minimal required core components of a modern computer system. It consists of three components as illustrated by Figure 2.1. The three components are the central processing unit, a cache, and the main memory, namely.




Figure 2.1: Hardware Model

### 2.1.1 Central Processing Unit

The *central processing unit (CPU)* is the core computational unit of computer systems. The purpose of a CPU is to process and execute a given program. A *program* consists of a sequence of instructions which operate on data. A program is processed sequentially instruction by instruction. Further, an *instruction* is a command with the purpose to perform some specific action, e.g., to add two numbers or to modify data. *Data* is the information required by a program for its execution, e.g., values for computations. The CPU consists of a limited number of so-called *registers*. A *register* is an extremely small and extremely fast memory unit which allows the CPU to execute computations. Registers are the only memory unit where the CPU is

able to apply arithmetic operations. The actual size of a single register and the number of registers available for computations depends on the architecture of the hardware.

For the purpose of this work only instructions reading or writing data are taken into account. To read data a so-called *load* instruction is required. And to write data a so-called *store* instructions is required. These two instructions access the data of a program stored at the main memory. E.g., for the purpose of computations which are done by the CPU it might be necessary to save a result for later computations. Such data can be write to main memory by executing a store instruction. However, there are many more instructions available on modern computer systems, e.g., arithmetical operations.

### 2.1.2 Main Memory

The *main memory* is a storage containing all data required to execute a program. Sometime the main memory is also called Random Access Memory (RAM). In general the CPU has to load and store data from the main memory to process a program. Furthermore, even the program itself is stored at the main memory while its execution. Each instruction of a program has to be loaded before the CPU is able to execute it. In case of an load instruction the CPU first loads the instruction itself. Second the CPU interprets the instructions and finds out that is has to execute a load of data. And third the CPU loads the actually required data into some of the available registers. Further, the main memory is structured in equally sized chunks of memory, e.g., 8 byte which is called represents the size of something called a *word*. Each such memory chunk can be accesses by the CPU via an unique identifier, its *physical address* or in short just *address*. If the CPU needs data for, e.g., a computation data has to be loaded via the load instruction `load &address`.

### 2.1.3 Cache

A *cache* is a small, high-speed memory which temporarily holds data of the addresses used by the currently processed program. The concept of caches has been introduced by Smith in his work [4]. Caches are based on two major observation. *Temporal locality*, if data is accessed it is likely that the same data is accessed again in the near future. *Spatial locality*, if data is accessed it is likely that other data near by is also accessed in the near future. Speaking about *accessing an address* is equivalent with *accessing data stored at an address in the main memory*. Same for cached addresses. For the CPU a cache is invisible. Independent if there is a cache or not the CPU always just wants to access a certain address. If there is cache present it simply takes less time to load the value of an address into the CPUs register. This is because caches are high-speed memory units. A *cache miss* occurs whenever the CPU wants to access an address which is not currently in the cache. Such a situation requires to load the data of the requested address from main memory into the cache. Before the value stored at this address is loaded into one of the registers. Such an operation is expensive as explained in section 2.5. If the requested address

is already in the cache it is not required to access the main memory, this case is called a *cache hit*. Since caches are small memory they are limited in the number of addresses which could be temporarily stored. In case the cache is full and a cache miss occurs it is required to make some space to load the requested address. The so-called *cache policy* decides which address has to be *evicted*, i.e., which address has to be written back into the main memory to get space. There are many different algorithms trying to make a good choice on the address to evict, e.g., least recently used (LRU). For more details see section 3.1

## 2.2 Memory Access Trace

The *memory access trace* represents all memory accesses for a given program. More precise the memory access trace is a sequence load and store instructions observed by analyzing a given program. Furthermore, for the purpose of this work it does not matter which value is stored at a certain address. For this reason the stored values are all dropped. This results in a sequence of tuples consisting of the instruction type, which is either *load* or *store*, and an address. Figure 2.4 shows an example of a memory access trace, addresses are annotated with a `&` known from languages like C.

Figure 2.2 shows a simple C program which is summing three numbers. At first all used variables are declared. Afterwards the variables `sum`, `x`, and `y` are initialized with the values 0, 1, and 2. Followed by the first computation, the value of `x` is added to `sum`'s value and stored in `sum`. Next the variable `z` is initialized with value 3. Then `sum` is increased by the value of `y`. Finally the computation is completed by adding the value of `z` to `sum`.

```
void main ()
{
    int sum, x, y, z;
    sum = 0;
    x = 1;
    y = 2;
    sum = sum + x;
    z = 3;
    sum = sum + y;
    sum = sum + z;
}
```

Figure 2.2: C code example of summing three numbers.

Figure 2.3 shows a snippet of assembly code generated by compiling the code of Figure 2.2. For compilation GCC 4.8.5 on Ubuntu 16.04 for AMD Opteron™ Processor 6376 with x86\_64 Architecture has been used. Since the declarations of the variables is only important for the compiler there no assembly code has been generated for these. For this reason the first line of assembly code shown in Figure 2.3 represents

the code generated for the C code `sum = 0;`. Other than expected the compiler has not generated a store operation instead the following code appears `movl $0, -16(%rbp)`.

```

movl    $0, -16(%rbp)
movl    $1, -12(%rbp)
movl    $2, -8(%rbp)
movl    -12(%rbp), %eax
addl    %eax, -16(%rbp)
movl    $3, -4(%rbp)
movl    -8(%rbp), %eax
addl    %eax, -16(%rbp)
movl    -4(%rbp), %eax
addl    %eax, -16(%rbp)

```

Figure 2.3: Assembly code snippet generated by compiling the code of Figure 2.2 with GCC 4.8.5 on Ubuntu 16.04.5 for AMD Opteron™ Processor 6376 with x86\_64 Architecture.

```

store &sum
store &x
store &y
load &sum
load &x
store &sum
store &z
load &sum
load &y
store &sum
load &sum
load &z
store &sum

```

Figure 2.4: Memory access trace of the assembly code shown in Figure 2.3.

Memory Access Type	Cost in Cycles
Cache load	1
Cache store	1
Memory load	5
Memory store	5

Table 2.1: Cost for memory access types

## 2.3 Liveness

## 2.4 Trace Transformation

## 2.5 Performance

Table 2.1 shows the costs for the different types of memory accesses. These numbers are taken from literature [1], <sup>1</sup>. The actual values are not that important than the relation of cache instruction costs to the memory instruction costs.

## 2.6 Problem Statement

Given a trace  $T$  of load and store instructions find metrics that characterize the trace performance for a given cache model  $C$  and implement an execution engine for computing their quantities.

---

<sup>1</sup><http://www.7-cpu.com/cpu/Skylake.html>



# Experimental Setup

## 3.1 Caches

### 3.1.1 Belady Cache

### 3.1.2 Belady Cache with Liveness Information

### 3.1.3 Least Recently Used Cache

### 3.1.4 Least Recently Used Cache with Liveness Information

## 3.2 Allocators

### 3.2.1 Original Allocator

### 3.2.2 Single Assignment Allocator

### 3.2.3 Compacting Allocator

#### 3.2.3.1 Compacting Allocator with Stack Semantic Free List

#### 3.2.3.2 Compacting Allocator with Queue Semantic Free List

#### 3.2.3.3 Compacting Allocator with Set Semantic Free List

## 3.3 Benchmarks

### 3.3.1 SPEC 2006 Benchmarks

This section describes the subset of benchmarks of the SPEC 2006 benchmark suite [2] which we used for our experiments. The benchmark descriptions below are taken from the SPEC 2006 paper.

**3.3.1.1 445.gobmk**

**Authors:** (in chronological order of contribution) are Man Lung Li, Wayne Iba, Daniel Bump, David Denholm, Gunnar Farnebäck, Nils Lohner, Jerome Dumonteil, Tommy Thorn, Nicklas Ekstrand, Inge Wallin, Thomas Traber, Douglas Ridgway, Teun Burgers, Tanguy Urvoy, Thien-Thi Nguyen, Heikki Levanto, Mark Vytlačil, Adriaan van Kessel, Wolfgang Manner, Jens Yllman, Don Dailey, Mans Ullerstam, Arend Bayer, Trevor Morris, Evan Berggren Daniel, Fernando Portela, Paul Pogonyshy, S.P. Lee, Stephane Nicolet and Martin Holters. General

**General Category:** Artificial intelligence - game playing.

**Description**<sup>1</sup>: The program plays Go and executes a set of commands to analyze Go positions.

**3.3.1.2 450.soplex**

**Authors:** Roland Wunderling, Thorsten Koch, Tobias Achterberg

**General Category:** Simplex Linear Program (LP) Solver

**Description:** 450.soplex is based on SoPlex Version 1.2.1. SoPlex solves a linear program using the Simplex algorithm. The LP is given as a sparse  $m$  by  $n$  matrix  $A$ , together with a right hand side vector  $b$  of dimension  $m$  and an objective function coefficient vector  $c$  of dimension  $n$ . The matrix is sparse in practice. SoPlex employs algorithms for sparse linear algebra, in particular a sparse LU-Factorization and solving routines for the resulting triangular equation systems.

**3.3.1.3 454.calculix**

**Author:** Guido D.C. Dhondt

**General Category:** Structural Mechanics

**Description**<sup>2</sup>: 454.calculix is based on CalculiX, a free software finite element code for linear and nonlinear three- dimensional structural applications. It uses classical theory of finite elements described in books such as [5]. CalculiX can solve problems such as static problems (bridge and building design), buckling, dynamic applications (crash, earthquake resistance) and eigenmode analysis (resonance phenomena).

**3.3.1.4 462 libquantum**

**Author:** Björn Butscher, Hendrik Weimer

**General Category:** Physics / Quantum Computing

---

<sup>1</sup>[www.gnu.org/software/gnugo/devel.html](http://www.gnu.org/software/gnugo/devel.html)

<sup>2</sup>[www.calculix.de](http://www.calculix.de)

**Description**<sup>3</sup>: libquantum is a library for the simulation of a quantum computer. Quantum computers are based on the principles of quantum mechanics and can solve certain computationally hard tasks in polynomial time. In 1994, Peter Shor discovered a polynomial-time algorithm for the factorization of numbers, a problem of particular interest for cryptanalysis, as the widely used RSA cryptosystem depends on prime factorization being a problem only to be solvable in exponential time. An implementation of Shor's factorization algorithm is included in libquantum. Libquantum provides a structure for representing a quantum register and some elementary gates. Measurements can be used to extract information from the system. Additionally, libquantum offers the simulation of decoherence, the most important obstacle in building practical quantum computers. It is thus not only possible to simulate any quantum algorithm, but also to develop quantum error correction algorithms. As libquantum allows to add new gates, it can easily be extended to fit the ongoing research, e.g. it has been deployed to analyze quantum cryptography.

#### 3.3.1.5 471 omnetpp

**Author:** Andras Varga, Omnest Global, Inc.

**General Category:** Discrete Event Simulation

**Description:** simulation of a large Ethernet network, based on the OMNeT++ discrete event simulation system<sup>4</sup>, using an ethernet model which is publicly available<sup>5</sup>. For the reference workload, the simulated network models a large Ethernet campus backbone, with several smaller LANs of various sizes hanging off each backbone switch. It contains about 8000 computers and 900 switches and hubs, including Gigabit Ethernet, 100Mb full duplex, 100Mb half duplex, 10Mb UTP, and 10Mb bus. The training workload models a small LAN. The model is accurate in that the CSMA/CD protocol of Ethernet and the Ethernet frame are faithfully modelled. The host model contains a traffic generator which implements a generic request-response based protocol. (Higher layer protocols are not modelled in detail.)

#### 3.3.1.6 483 xalancbmk

**Author:** IBM Corporation, Apache Inc, plus modifications for SPEC purposes by Christopher Cambly, Andrew Godbout, Neil Graham, Sasha Kasapinovic, Jim McInnes, June Ng, Michael Wong. Primary contact: Michael Wong

---

<sup>3</sup><http://www.libquantum.de>

<sup>4</sup>[www.omnetpp.org](http://www.omnetpp.org)

<sup>5</sup><http://ctieware.eng.monash.edu.au/twiki/bin/view/Simulation/EtherNet>

**General Category:** XSLT processor for transforming XML documents into HTML, text, or other XML document types

**Description:** a modified version of Xalan-C++<sup>6</sup>, an XSLT processor written in a portable subset of C++ . Xalan-C++ version 1.8 is a robust implementation of the W3C Recommendations for XSL Transformations (XSLT)<sup>7</sup> and the XML Path Language (XPath)<sup>8</sup>. It works with a compatible release of the Xerces-C++<sup>9</sup> XML parser: Xerces-C++ version 2.5.0. The XSLT language is used to compose XSL stylesheets. An XSL stylesheet contains instructions for transforming XML documents from one document type to another document type (XML, HTML, or other). In structural terms, an XSL stylesheet specifies the transformation of one tree of nodes (the XML input) into another tree of nodes (the output or transformation result). Modifications for SPEC benchmarking purposes include: combining code to make a standalone executable, removing compiler incompatibilities and improving standard conformance, changing output to display intermediate values, removing large parts of unexecuted code, and moving all the include locations to fit better into the SPEC harness.

### 3.3.2 Google JavaScript Engine (V8) Benchmarks

This section describes the subset of benchmarks of the Octane benchmark suite [?] which we used for our experiments. The benchmark descriptions below are taken from the Octane website.

#### 3.3.2.1 Richards

OS kernel simulation benchmark, originally written in BCPL by Martin Richards<sup>10</sup> (539 lines).

**Main focus:** property load/store, function/method calls

**Secondary focus:** code optimization, elimination of redundant code

#### 3.3.2.2 Raytrace

Ray tracer benchmark based on code by Adam Burmister<sup>11</sup> (904 lines).

**Main focus:** argument object, apply

**Secondary focus:** prototype library object, creation pattern

---

<sup>6</sup><http://xml.apache.org/xalan-c/>

<sup>7</sup><http://www.w3.org/TR/xslt>

<sup>8</sup><http://www.w3.org/TR/xpath>

<sup>9</sup><http://xml.apache.org/xerces-c>

<sup>10</sup><http://www.cl.cam.ac.uk/~mr10/>

<sup>11</sup><http://burmister.com>

### 3.3.2.3 Deltablue

One-way constraint solver<sup>12</sup>, originally written in Smalltalk by John Maloney and Mario Wolczko (880 lines).

**Main focus:** polymorphism

**Secondary focus:** OO-style programming

## 3.4 Metrics

### 3.4.1 Address Access

### 3.4.2 Access Distance

### 3.4.3 Live Addresses

### 3.4.4 Liveness Length

## 3.5 Trace generation (Valgrind / Cachegrind)

---

<sup>12</sup><http://constraints.cs.washington.edu/deltablue/>


# CHAPTER 4

## Experiments

4.1 8 Byte Cache Line Size

4.2 64 Byte Cache Line Size

CHAPTER 5




## Related Work

# CHAPTER 6

## Conclusion



# CHAPTER 7



## Future Work

## Acronyms

<b>CPA</b>	cycles per access
<b>CPU</b>	central processing unit
<b>LRU</b>	least recently used
<b>OS</b>	operating system
<b>SATrace</b>	single assignment trace
<b>SSA</b>	static single assignment
<b>VM</b>	virtual machine
<b>RAM</b>	Random Access Memory

# Appendix

# List of Figures

2.1	Hardware Model . . . . .	2
2.2	C code example of summing three numbers. . . . .	4
2.3	Assembly code snippet generated by compiling the code of Figure 2.2 with GCC 4.8.5 on Ubuntu 16.04.5 for AMD Opteron™ Processor 6376 with x86_64 Architecture. . . . .	5
2.4	Memory access trace of the assembly code shown in Figure 2.3. . . . .	5

# List of Tables

2.1	Cost for memory access types . . . . .	6
-----	--	---

# List of Algorithms

# Acronyms

**RDBMS** Relational Database Management System

**SQL** Structured Query Language

# Bibliography

- [1] Ulrich Drepper. What every programmer should know about memory. *Red Hat, Inc*, 11:2007, 2007.
- [2] John L Henning. Spec cpu2006 benchmark descriptions. *ACM SIGARCH Computer Architecture News*, 34(4):1–17, 2006.
- [3] Bruce Jacob, Spencer Ng, and David Wang. *Memory systems: cache, DRAM, disk*. Morgan Kaufmann, 2010.
- [4] Alan Jay Smith. Cache memories. *ACM Computing Surveys (CSUR)*, 14(3):473–530, 1982.
- [5] Olgierd Cecil Zienkiewicz, Robert Leroy Taylor, Olgierd Cecil Zienkiewicz, and Robert Lee Taylor. *The finite element method*, volume 3. McGraw-hill London, 1977.

**TODO:** dummy cite to activate bib: [3]